

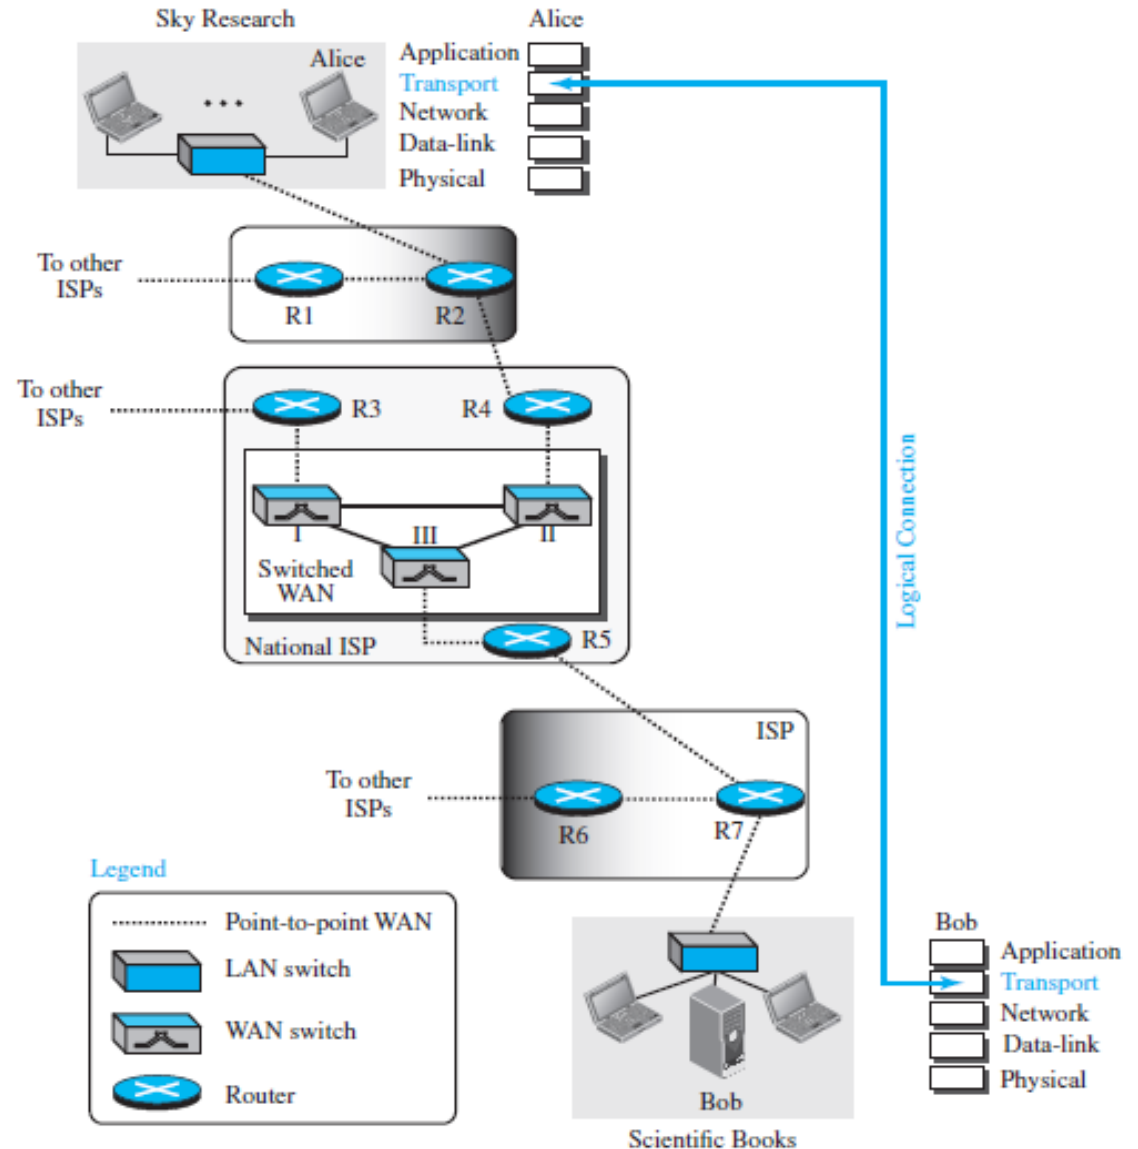
Transport Layer

Dr. Subrat K. Dash

Transport Layer – The Link

- Application layer
 - Communication for specific applications
 - E.g., HyperText Transfer Protocol (HTTP), File Transfer Protocol (FTP), Network News Transfer Protocol (NNTP)
- Transport layer
 - Communication between processes (e.g., socket)
 - Relies on network layer and serves the application layer
 - E.g., TCP and UDP
- Network layer
 - Logical communication between nodes
 - Hides details of the link technology
 - E.g., IP

Transport Layer




Transport Protocols

- Run on end hosts
 - Sender: breaks application messages into segments,
and passes to network layer
 - Receiver: reassembles segments into messages,
passes to application layer
- Multiple transport protocol available to applications
 - Internet: TCP and UDP

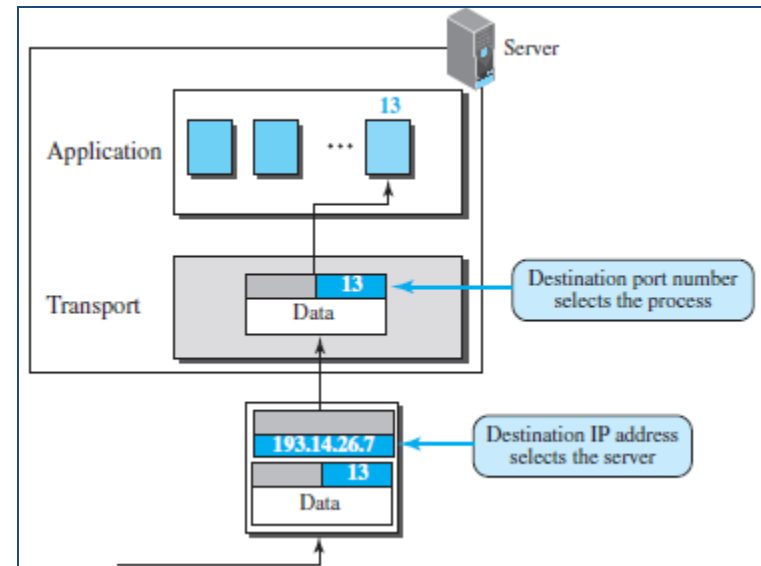

Transport Layer Services

- Process-to-Process Communication




Transport Layer Services

- Addressing: Port Numbers (16 bits)
 - Client-server paradigm
 - Port numbers
 - Client: Ephemeral port (>1023)
 - Server: Well-known port



ICANN Ranges



```
$grep tftp/etc/services
```

```
tftp 69/tcp
```

```
tftp 69/udp
```

```
$grep snmp/etc/services
```

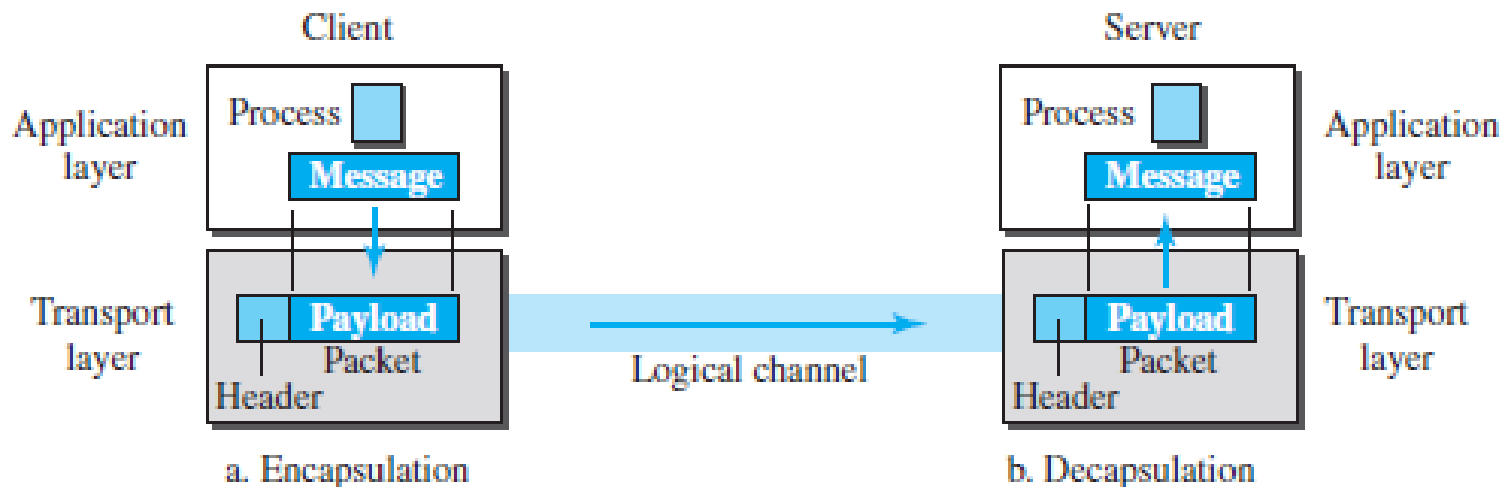
```
snmp161/tcp#Simple Net Mgmt Proto
```

```
snmp161/udp#Simple Net Mgmt Proto
```

```
snmptrap162/udp#Traps for SNMP
```

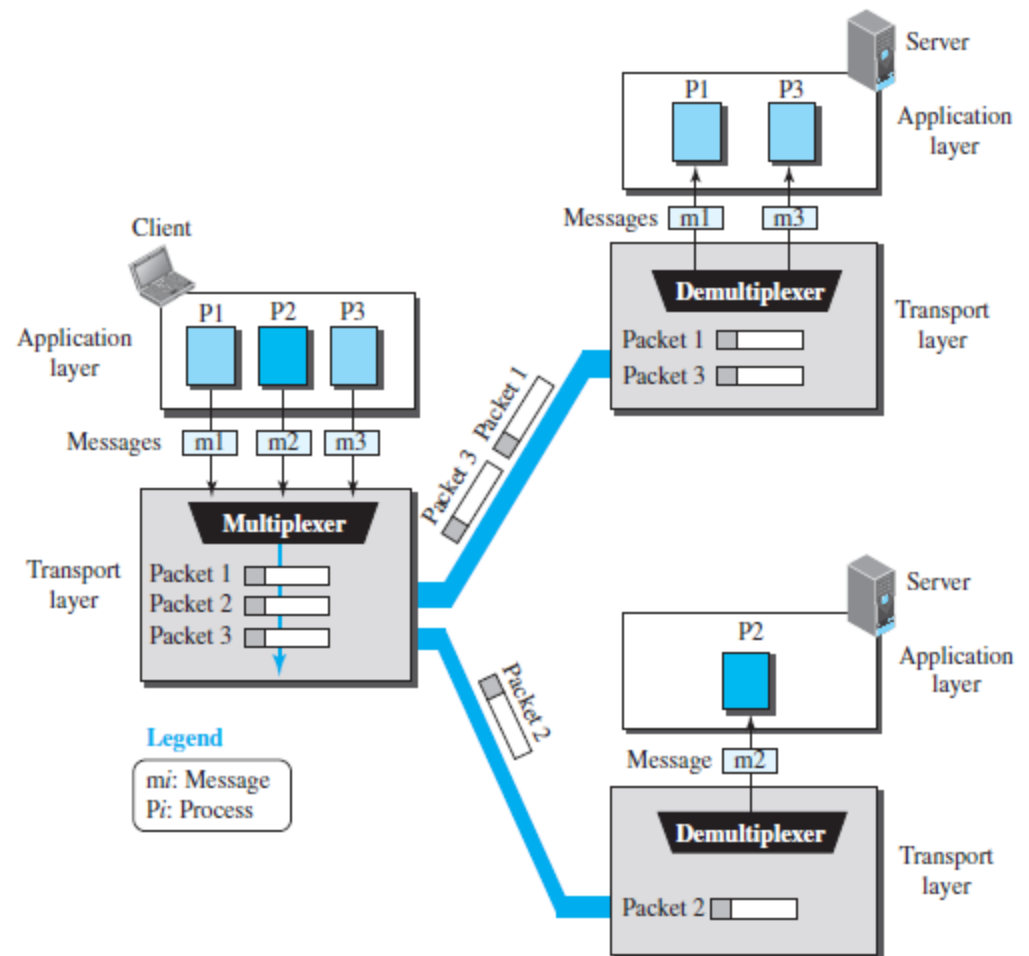
Transport Layer Services

- Encapsulation and Decapsulation
 - User datagrams, Segments, Packets



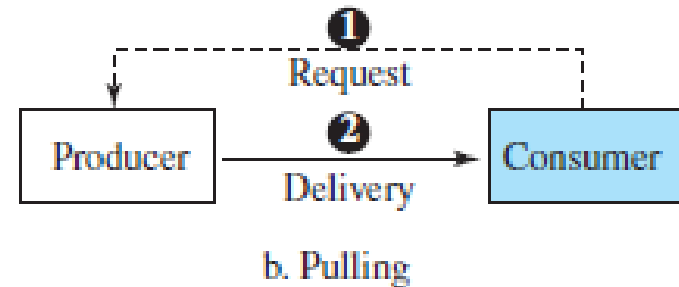

Transport Layer Services

- Multiplexing and Demultiplexing



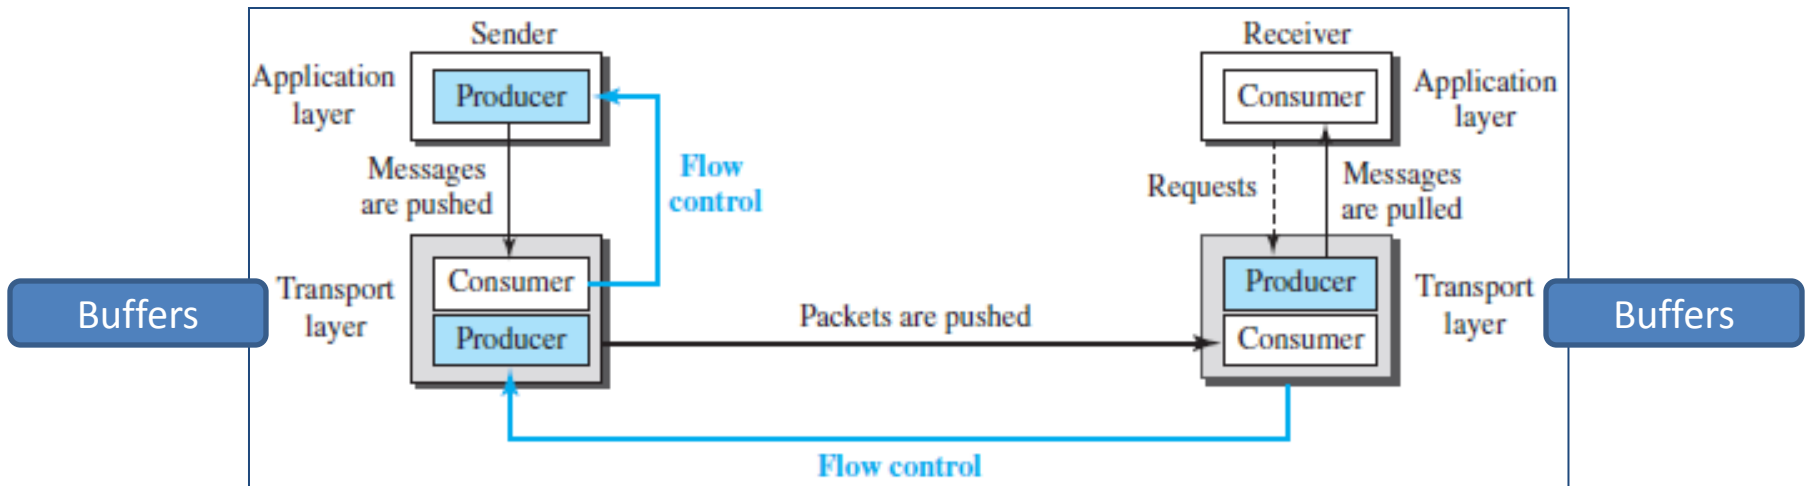

Transport Layer Services

- Flow Control
 - Items are produced faster than they can be consumed.




Transport Layer Services

- Flow Control
 - Items are produced faster than they can be consumed.



Transport Layer Services

- Error Control
 - Detecting and discarding corrupted packets.
 - Keeping track of lost and discarded packets and resending them.
 - Recognizing duplicate packets and discarding them.
 - Buffering out-of-order packets until the missing packets arrive.



Transport Layer Services

- Error Control
 - Sequence Number
 - Acknowledgement



a. Four packets have been sent.



b. Five packets have been sent.



c. Seven packets have been sent;
window is full.



d. Packet 0 has been acknowledged;
window slides.

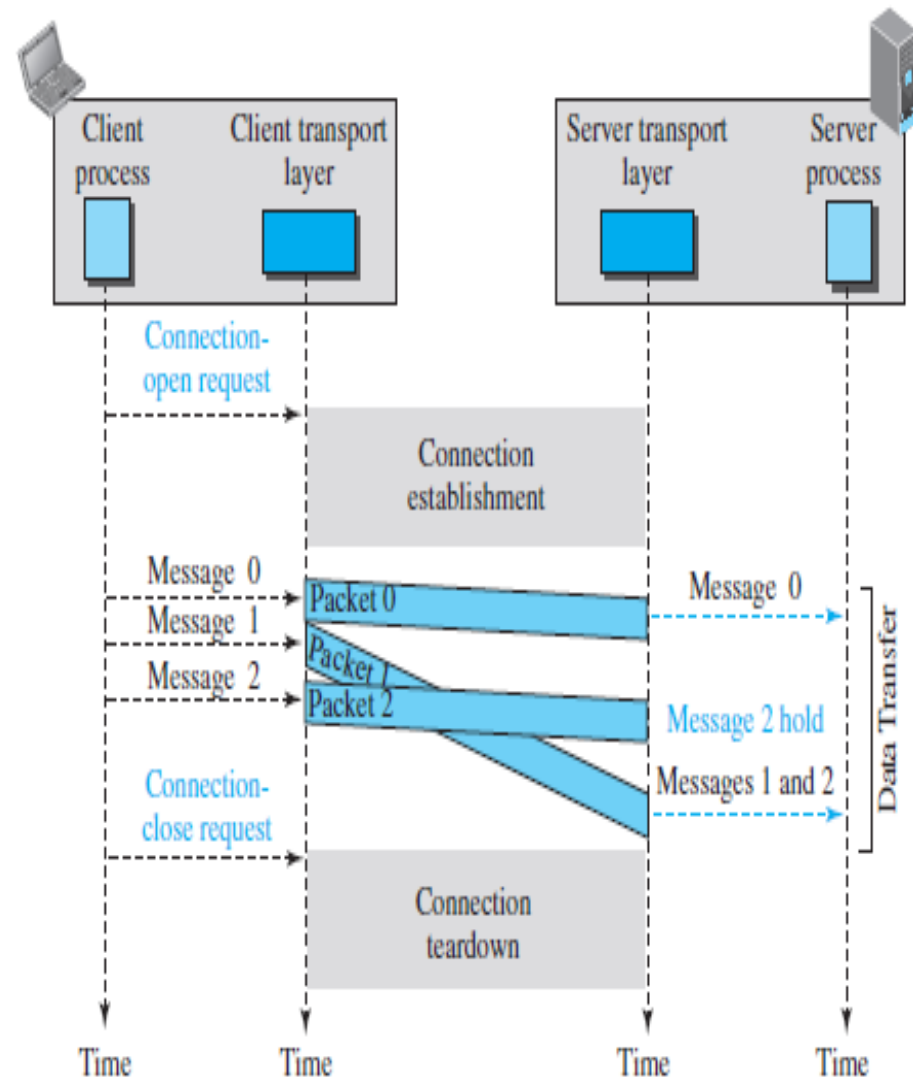
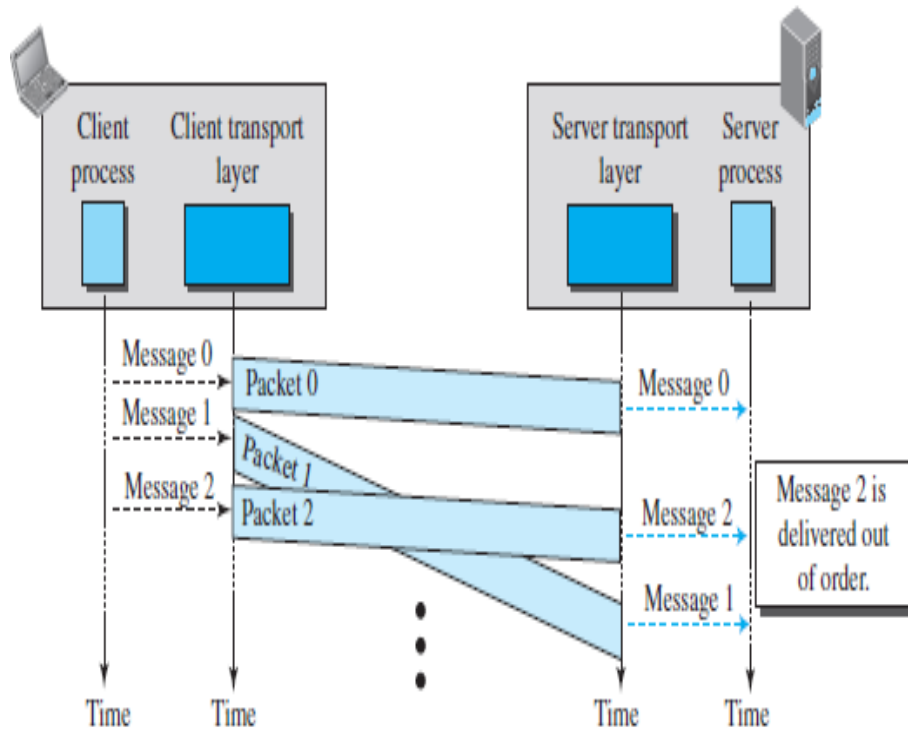
Transport Layer Services

- Congestion Control
 - Congestion
 - Results due to imbalance between **load** of network and **capacity** of network
 - Waiting
 - TCP

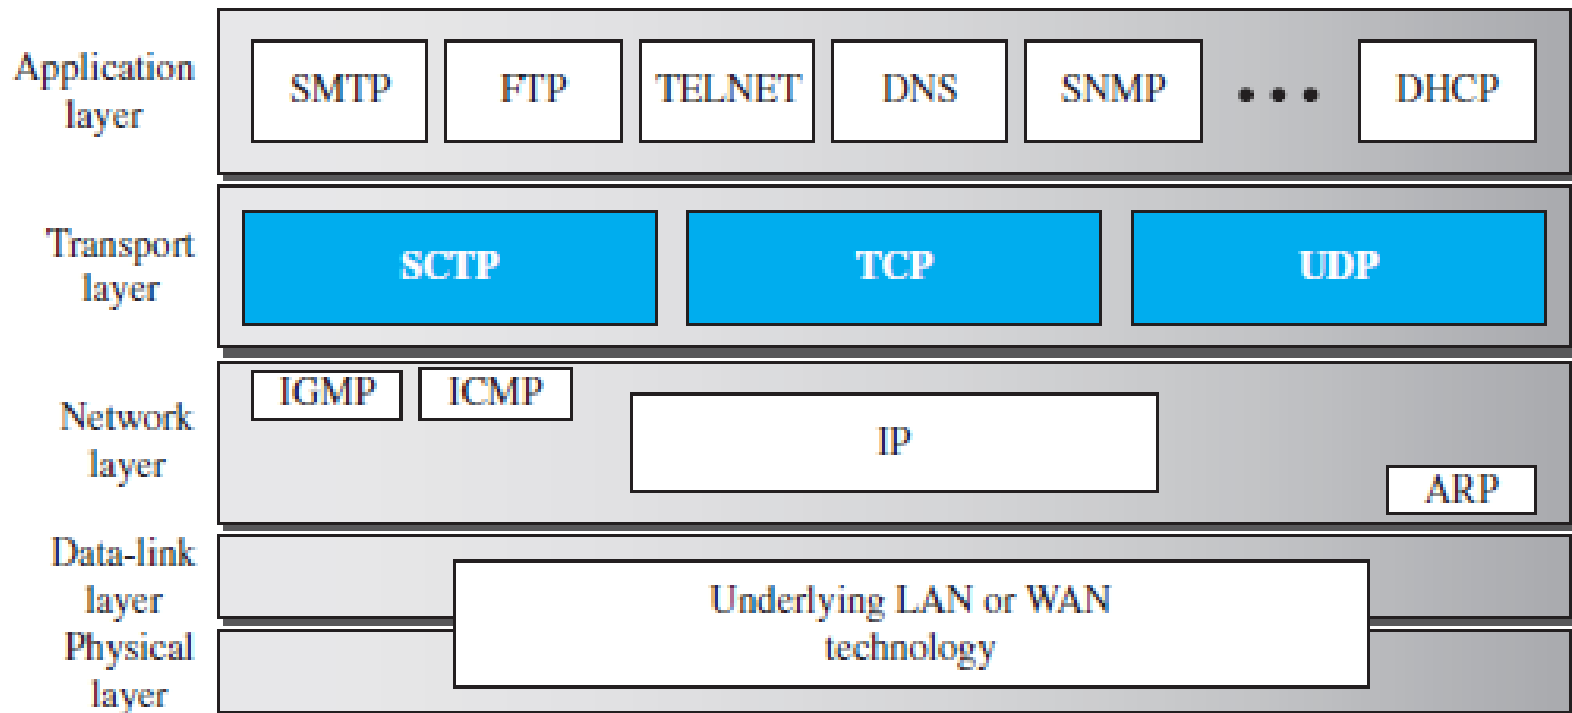
Protocol Types

- Connectionless
 - Independency between packets (TL)
 - Different paths for different datagrams belonging to the same message (NL)
- Connection-oriented
 - Dependency between packets (TL)
 - Same path for different datagrams belonging to the same message (NL)

Connectionless vs. Connection-oriented



Transport-Layer Protocols in TCP/IP Protocol Suite




Sample list of well-known ports


<i>Port</i>	<i>Protocol</i>	<i>UDP</i>	<i>TCP</i>	<i>SCTP</i>	<i>Description</i>
7	Echo	√	√	√	Echoes back a received datagram
9	Discard	√	√	√	Discards any datagram that is received
11	Users	√	√	√	Active users
13	Daytime	√	√	√	Returns the date and the time
17	Quote	√	√	√	Returns a quote of the day
19	Chargen	√	√	√	Returns a string of characters
20	FTP-data		√	√	File Transfer Protocol
21	FTP-21		√	√	File Transfer Protocol
23	TELNET		√	√	Terminal Network
25	SMTP		√	√	Simple Mail Transfer Protocol
53	DNS	√	√	√	Domain Name Service
67	DHCP	√	√	√	Dynamic Host Configuration Protocol
69	TFTP	√	√	√	Trivial File Transfer Protocol
80	HTTP		√	√	HyperText Transfer Protocol
111	RPC	√	√	√	Remote Procedure Call
123	NTP	√	√	√	Network Time Protocol
161	SNMP-server	√			Simple Network Management Protocol
162	SNMP-client	√			Simple Network Management Protocol

User Datagram Protocol (UDP)

- Connectionless, unreliable TL Protocol
- User datagram packet format



a. UDP user datagram




b. Header format


UDP Services

- Process-to-process communication
- Encapsulation and decapsulation
- Multiplexing and demultiplexing
- Flow control
- Error control
- Congestion control
- Connectionless services
- Checksum

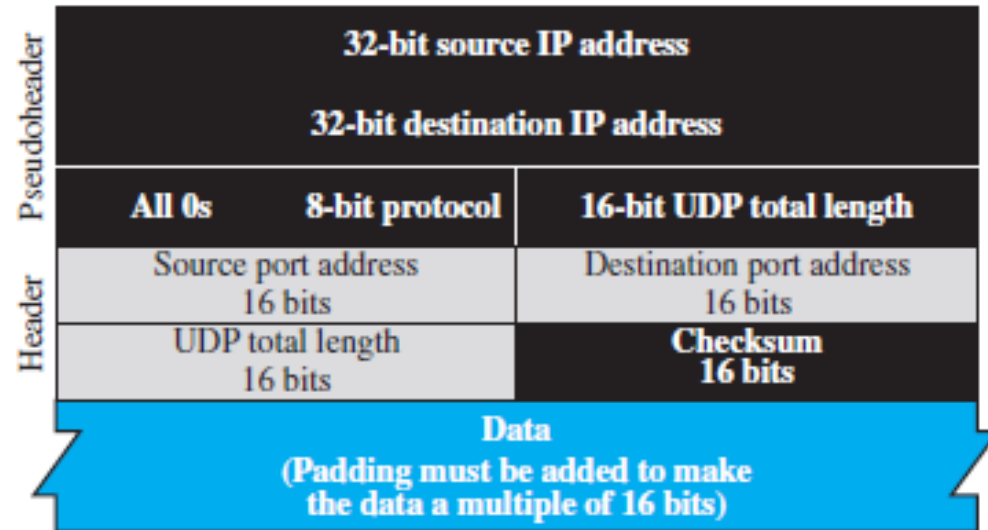
UDP: Checksum



a. UDP user datagram



b. Header format



TCP

- Connection oriented
 - Explicit set-up and tear-down of TCP session
- Stream-of-bytes service
 - Sends and receives a stream of bytes, not messages
- Reliable, in-order delivery
 - Checksums to detect corrupted data
 - Acknowledgments & retransmissions for reliable delivery
 - Sequence numbers to detect losses and reorder data
- Flow control
 - Prevent overflow of the receiver's buffer space
- Congestion control
 - Adapt to network congestion for the greater good

An Analogy: Talking on a Cell Phone

- Alice and Bob on their cell phones
 - Both Alice and Bob are talking
- What if Alice couldn't understand Bob?
 - Bob asks Alice to repeat what she said
- What if Bob hasn't heard Alice for a while?
 - Is Alice just being quiet?
 - Or, have Bob and Alice lost reception?
 - How long should Bob just keep on talking?
 - Maybe Alice should periodically say “uh huh”
 - ... or Bob should ask “Can you hear me now?”

Some Take-Aways from the Example

- Acknowledgments from receiver
 - Positive: “okay” or “ACK”
 - Negative: “please repeat that” or “NACK”
- Timeout by the sender (“stop and wait”)
 - Don’t wait indefinitely without receiving some response
 - ... whether a positive or a negative acknowledgment
- Retransmission by the sender
 - After receiving a “NACK” from the receiver
 - After receiving no feedback from the receiver

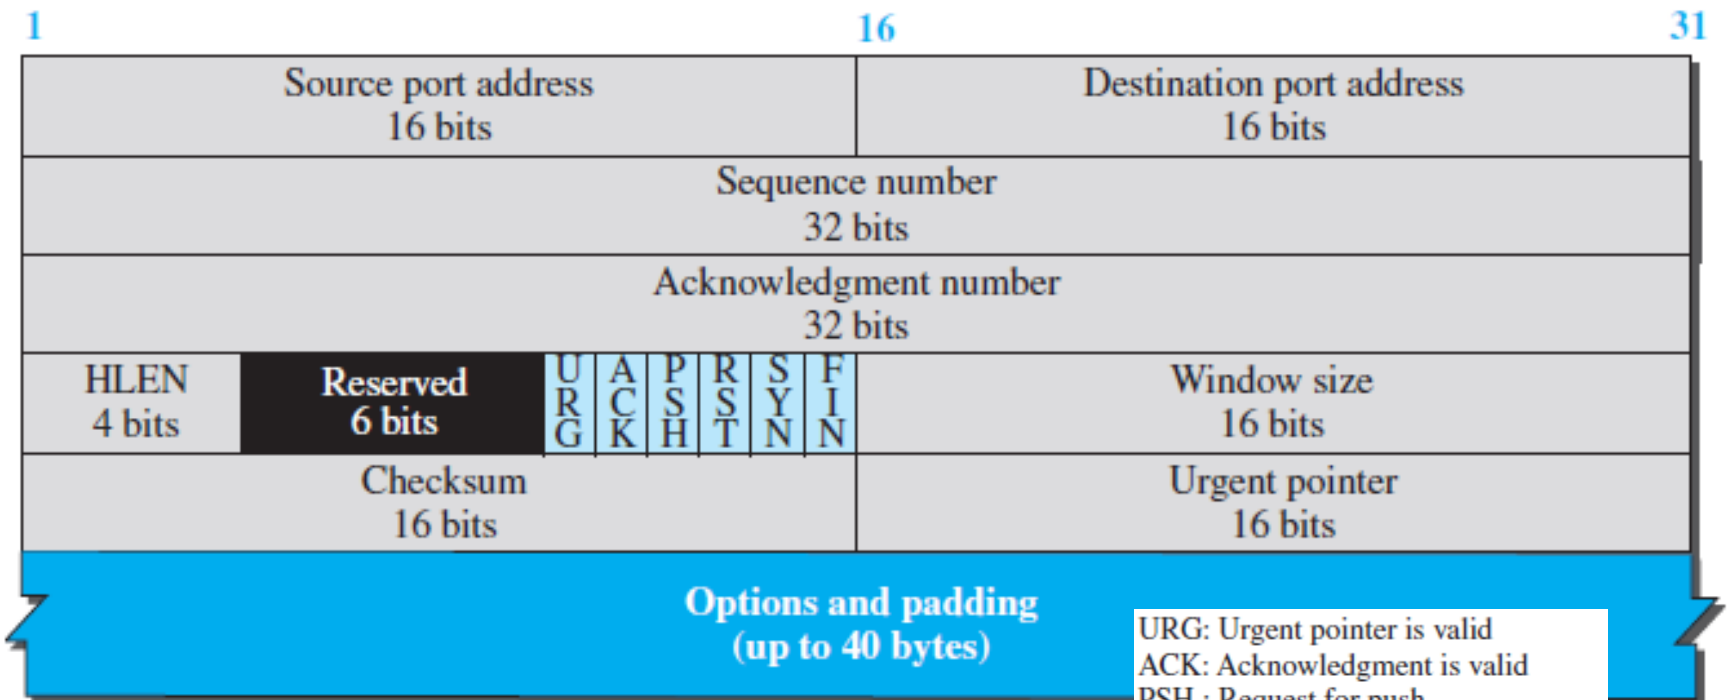
TCP Services

- Process-to-process communication
- Stream delivery service
- Full-Duplex communication
- Multiplexing and demultiplexing
- Connection-oriented service
- Reliable service

TCP Segment



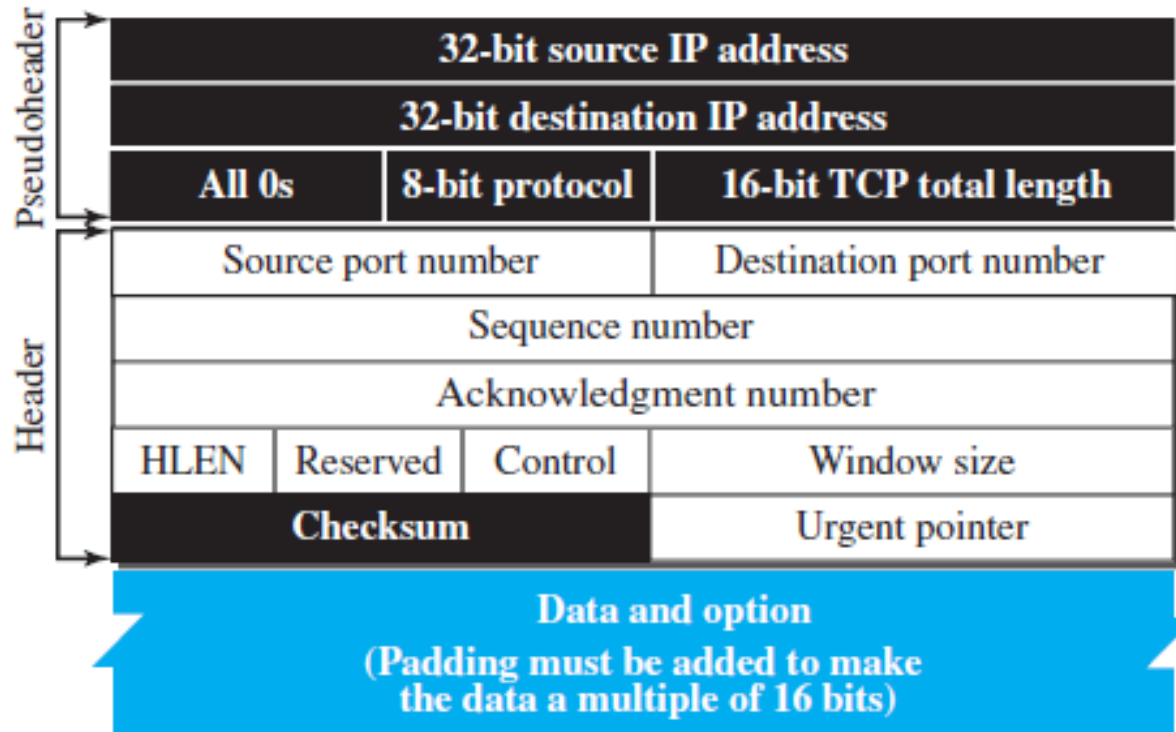
a. Segment



b. Header

URG: Urgent pointer is valid
 ACK: Acknowledgment is valid
 PSH: Request for push
 RST: Reset the connection
 SYN: Synchronize sequence numbers
 FIN: Terminate the connection

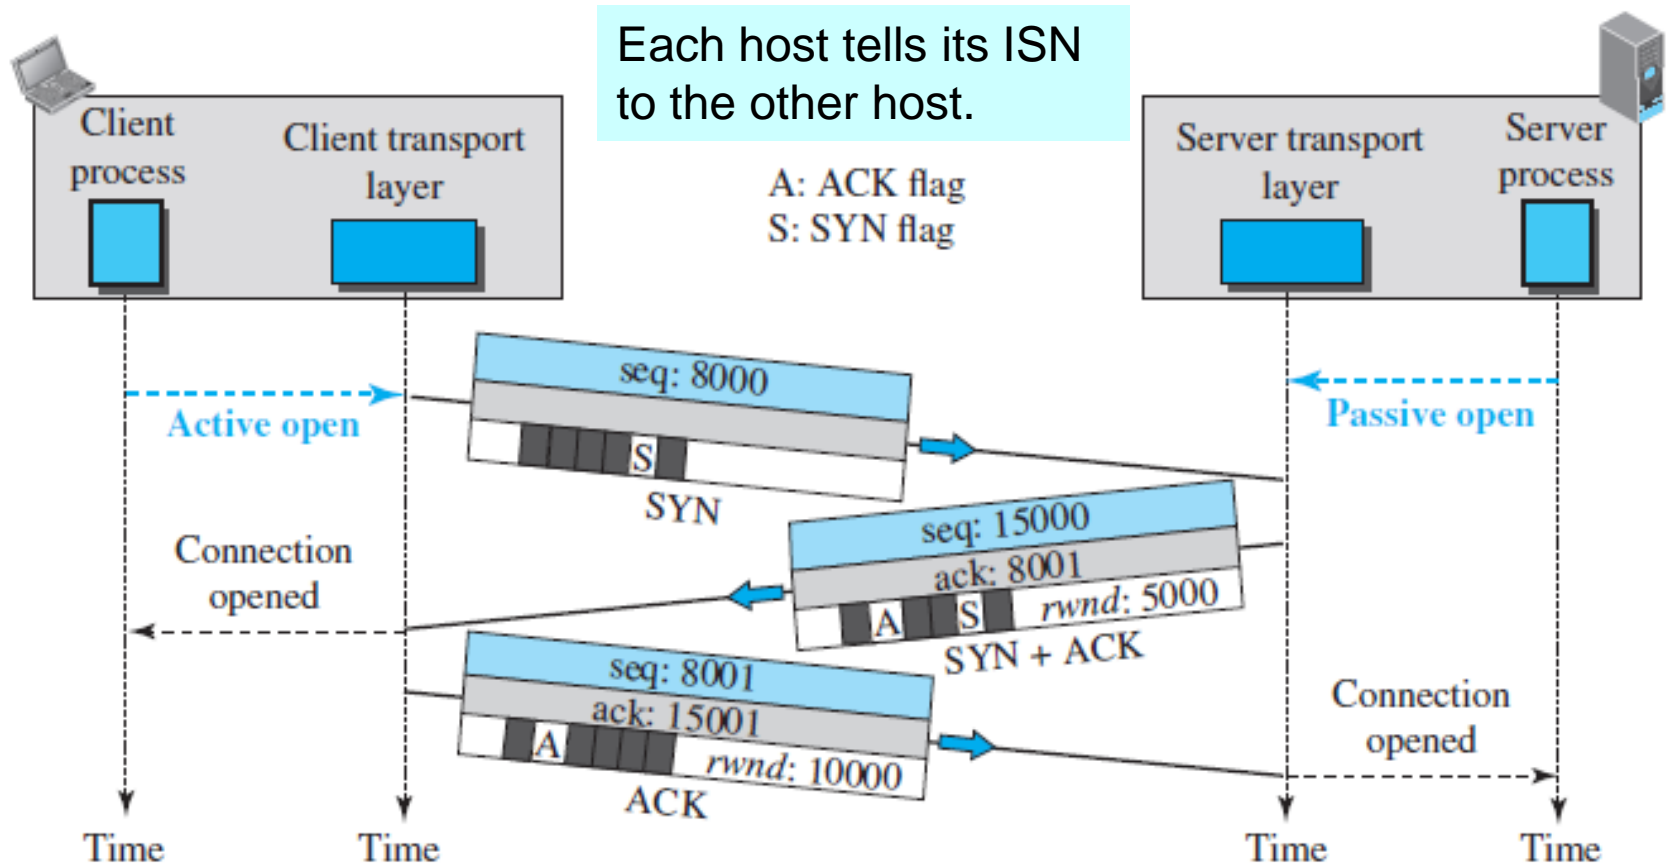
TCP – Checksum



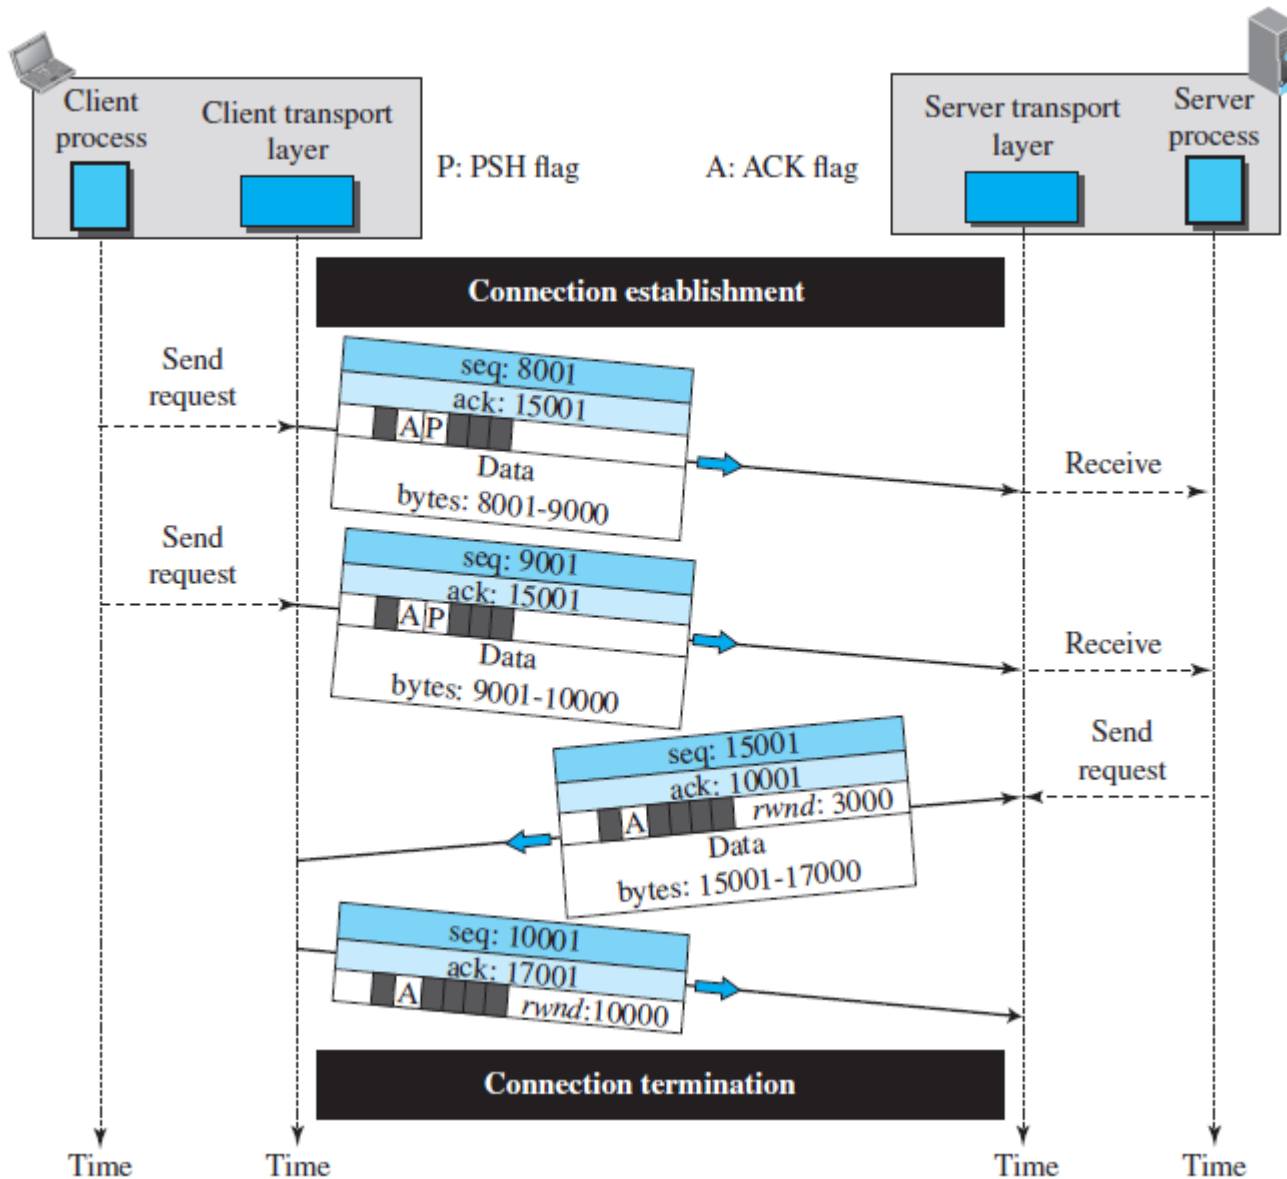
The use of the checksum in TCP is mandatory.

TCP Connection

- Connection establishment
 - Three-way handshaking

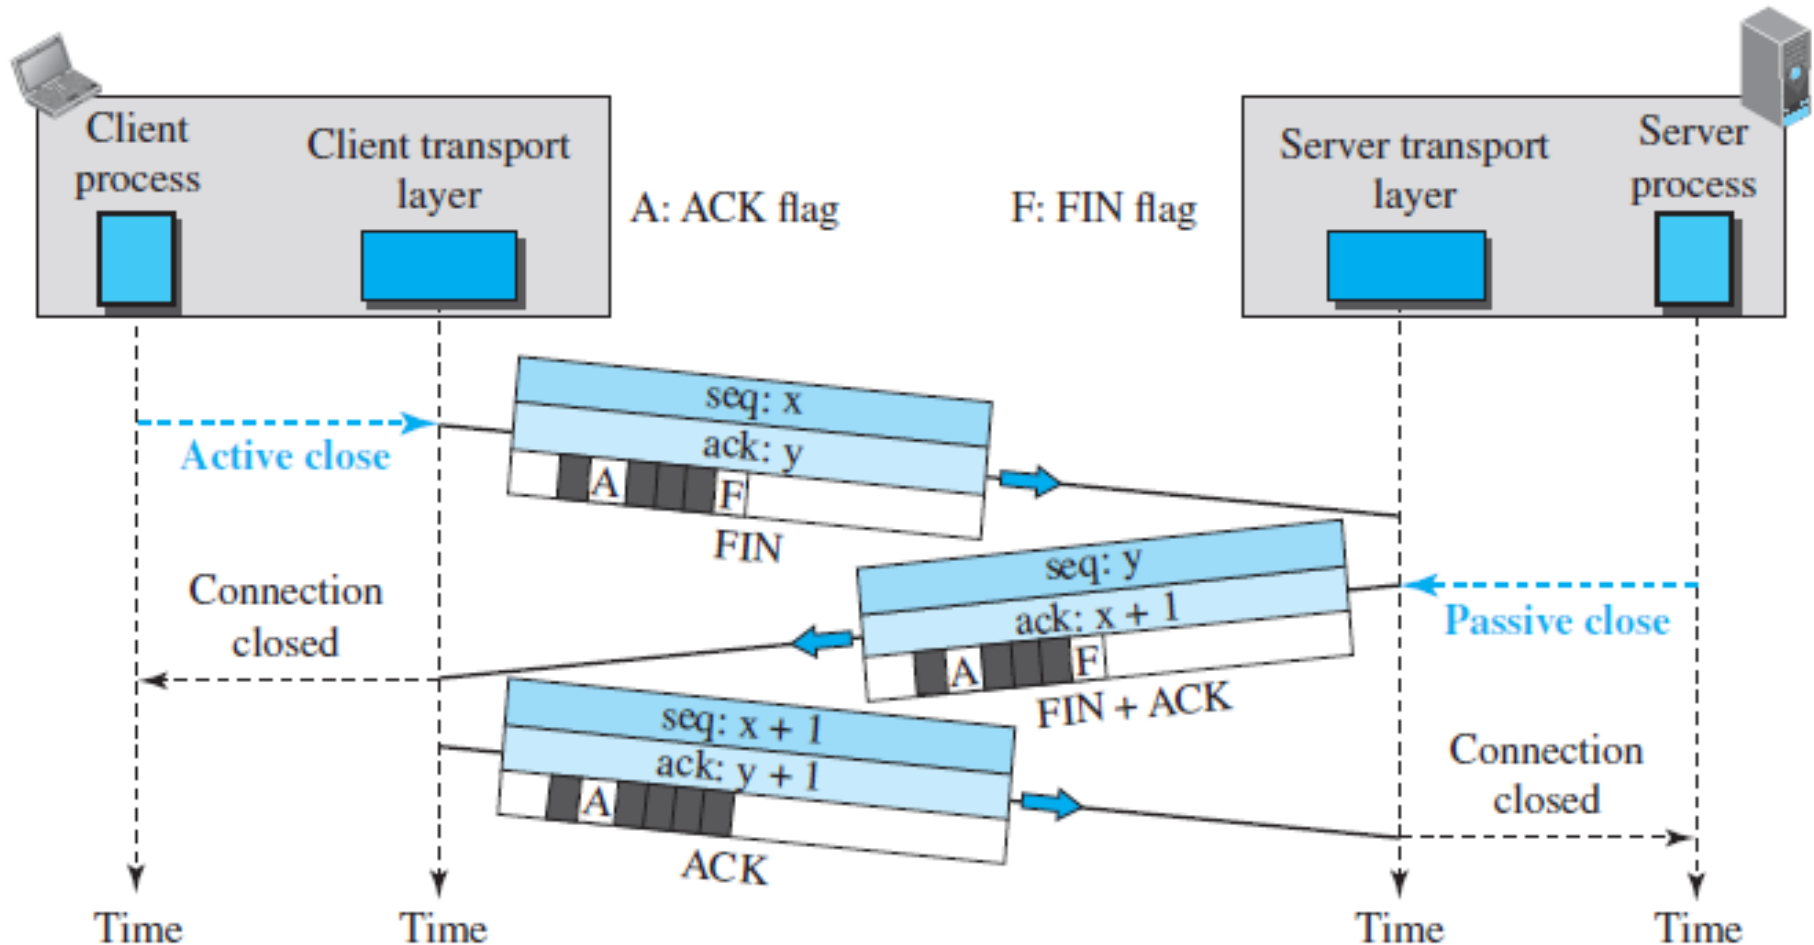


Data Transfer

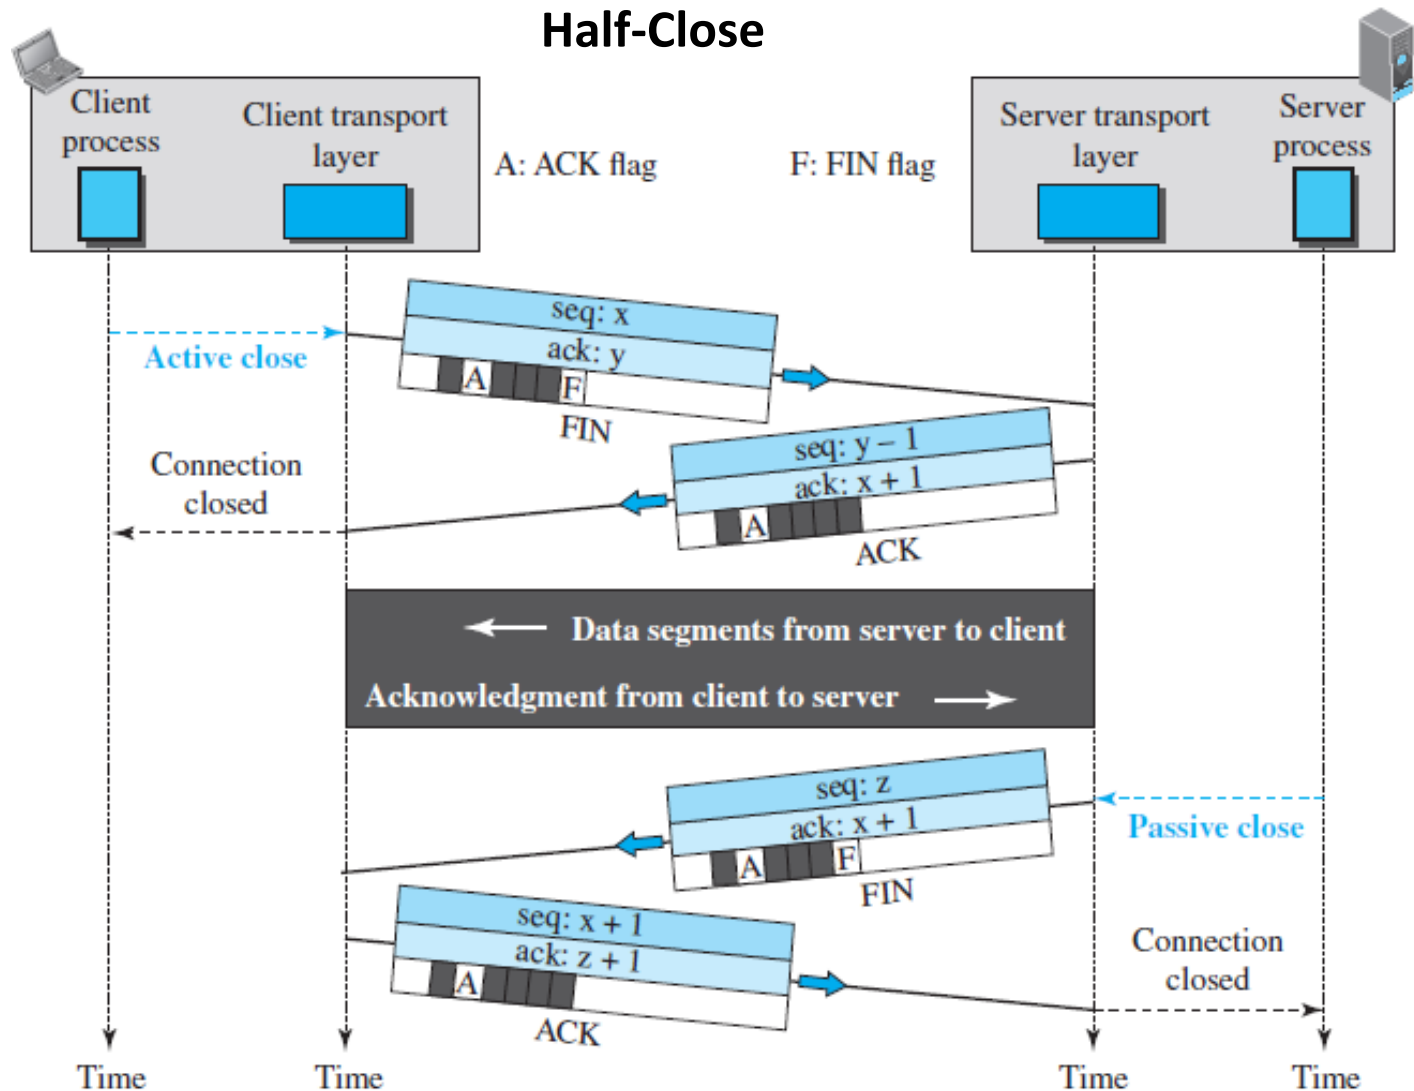


Connection Termination


Three-Way Handshaking



Connection Termination

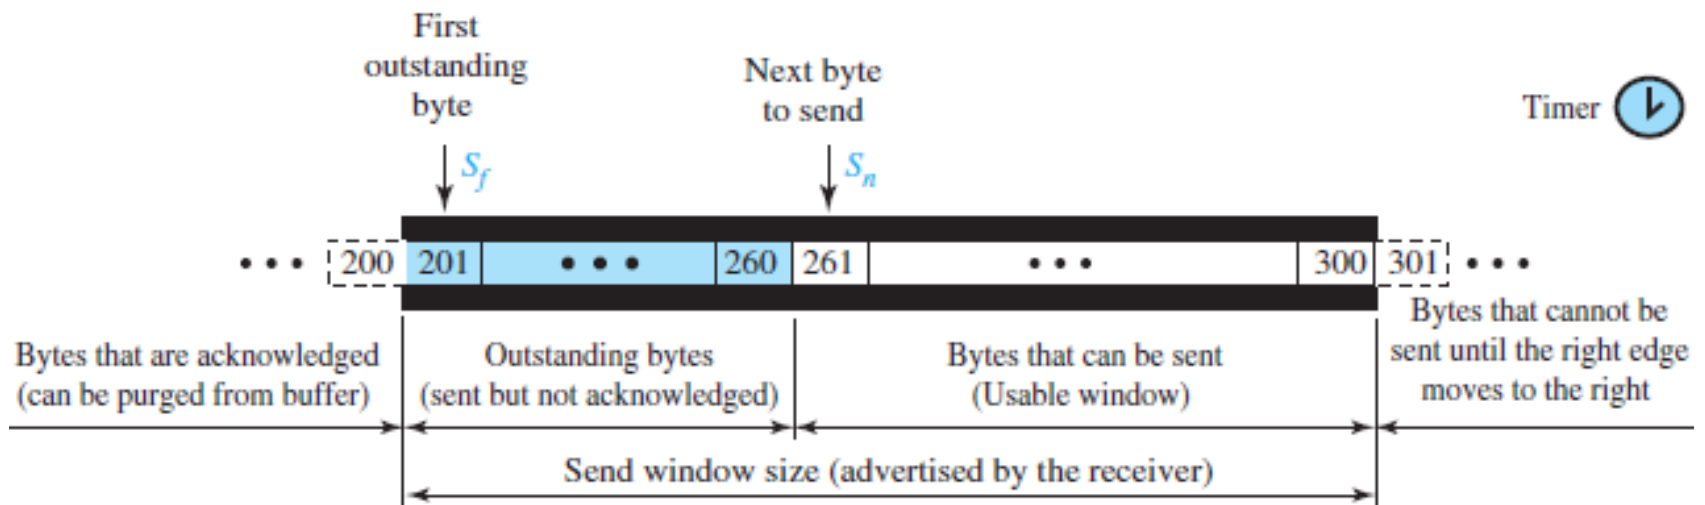


Time-line diagram for a common scenario



Windows in TCP

Send Window



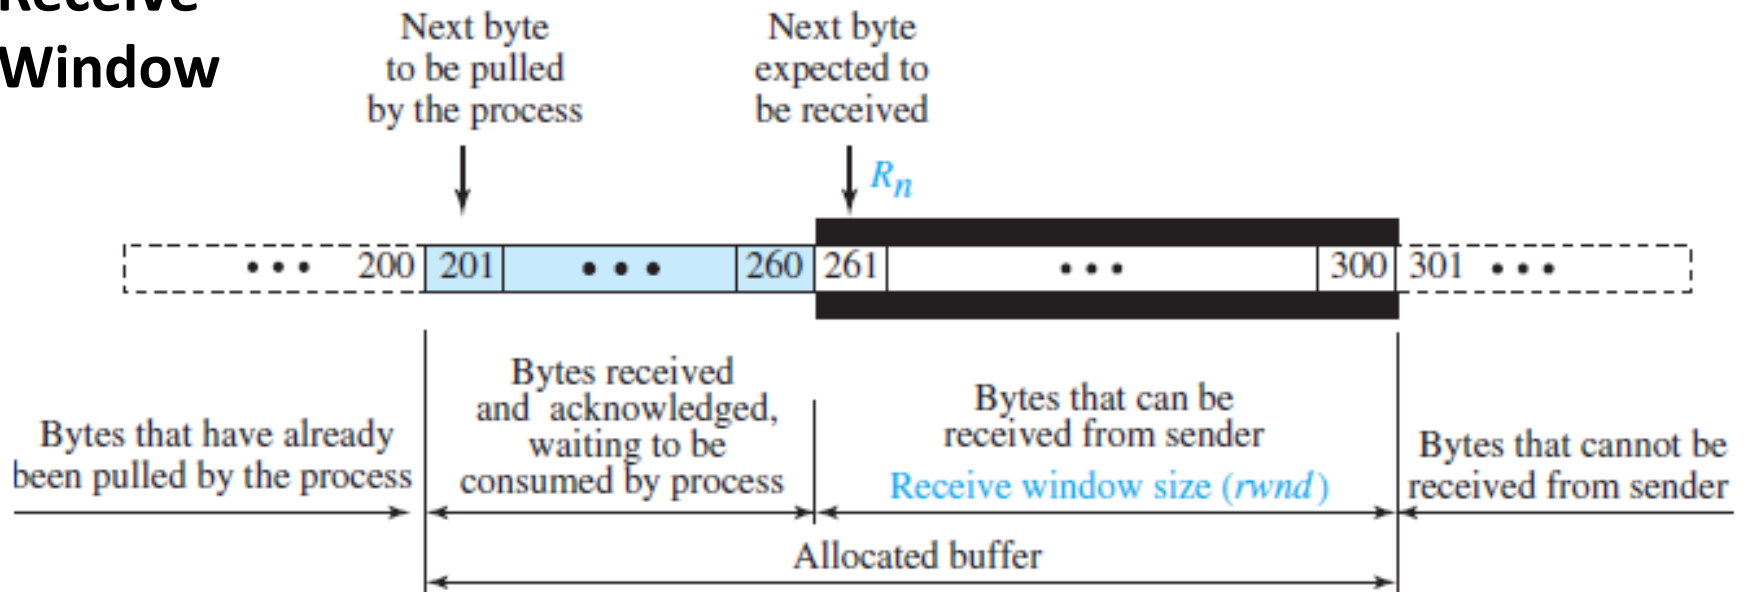
a. Send window




b. Opening, closing, and shrinking send window

Windows in TCP

Receive Window

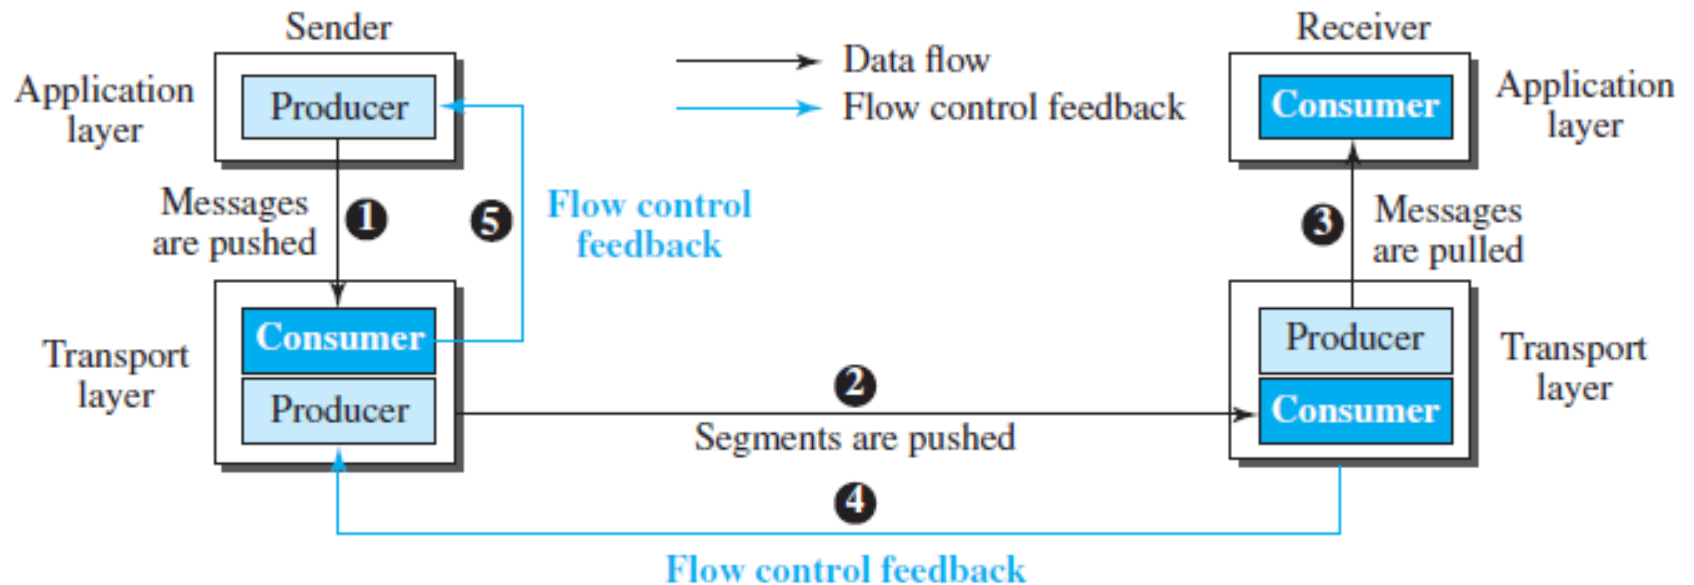


a. Receive window and allocated buffer

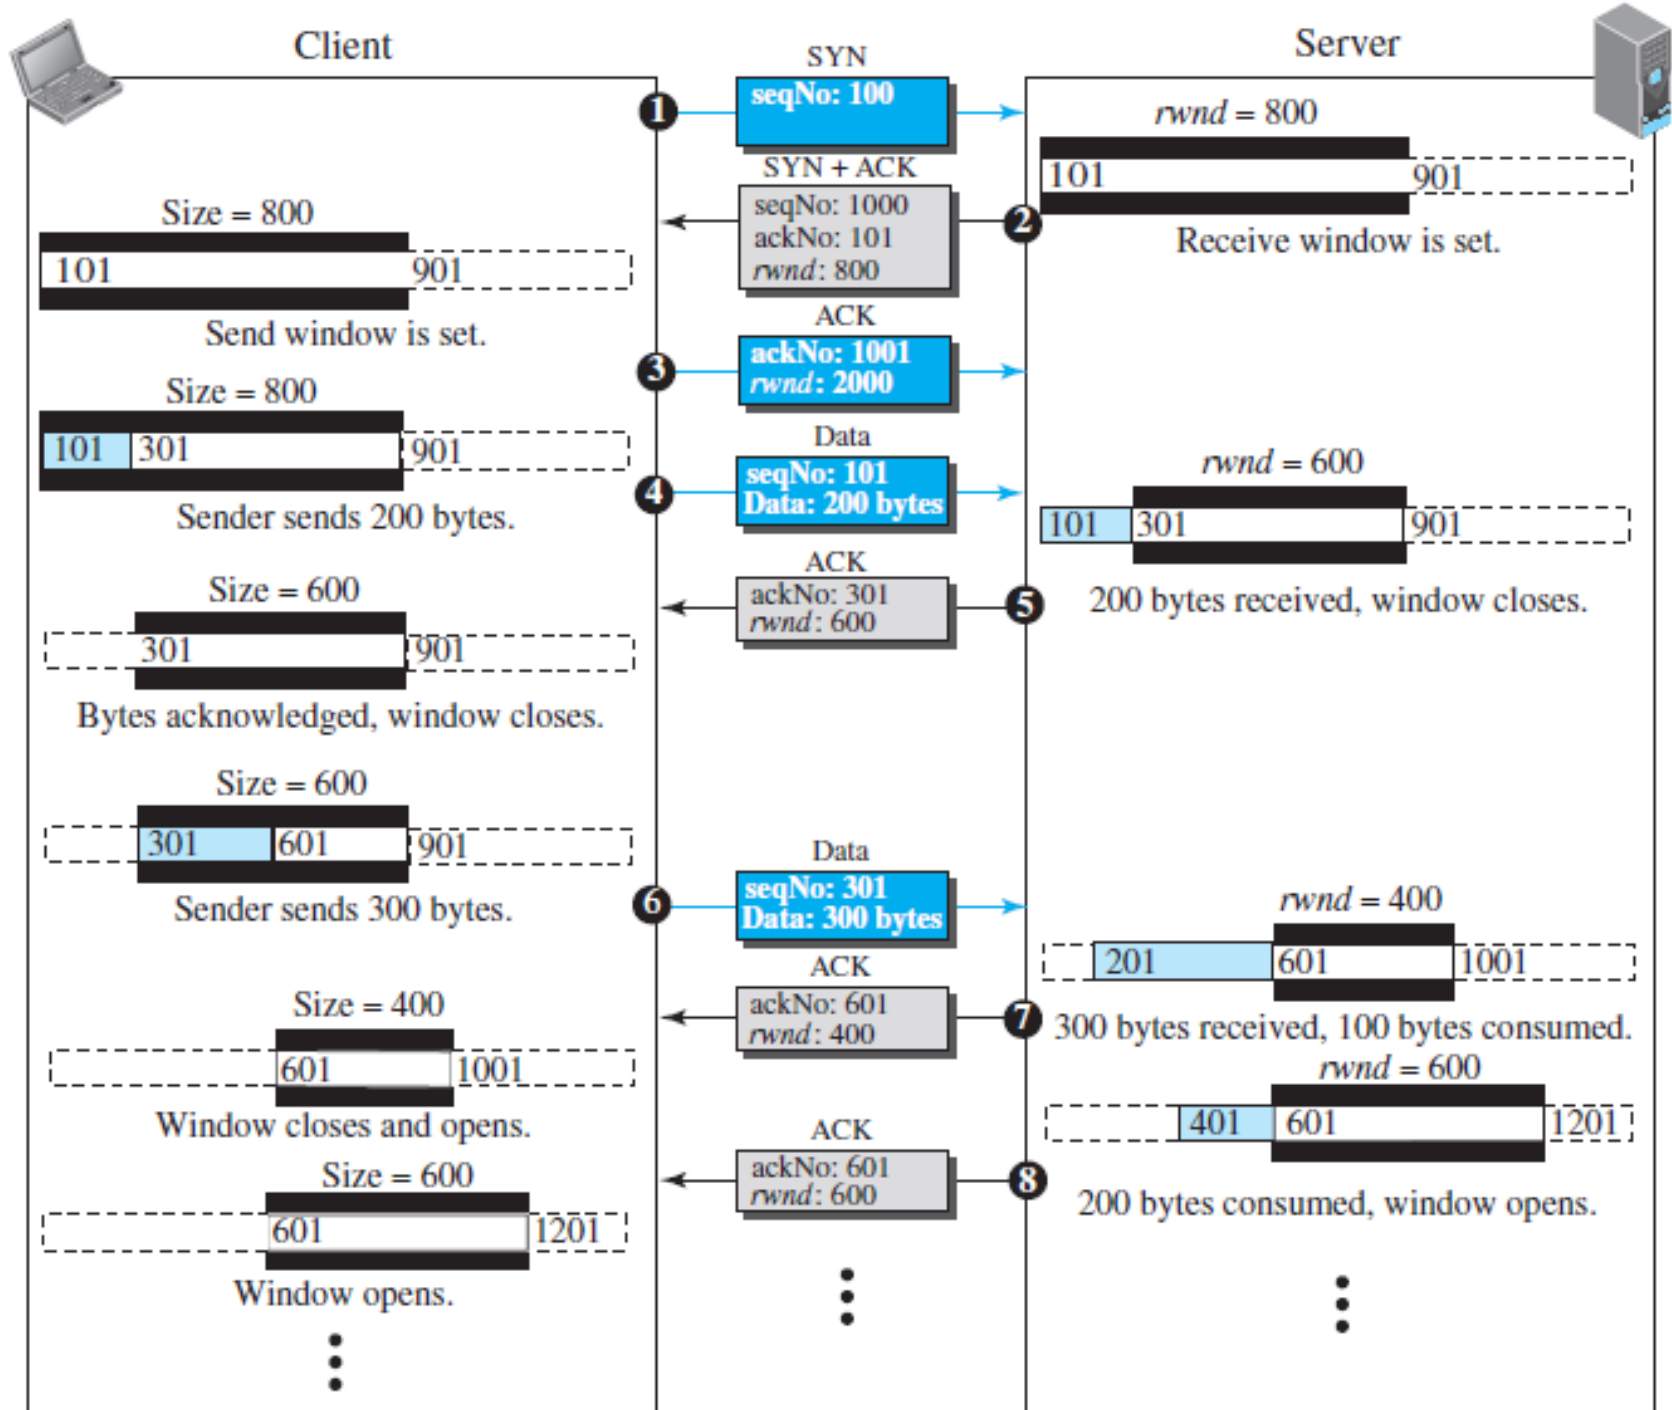


b. Opening and closing of receive window

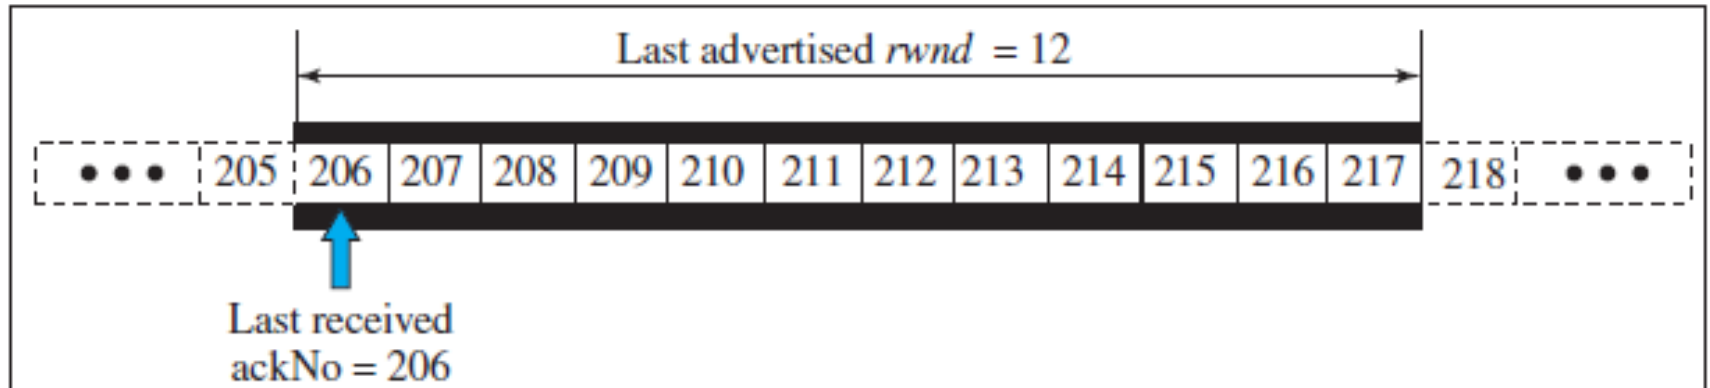
TCP – Flow Control



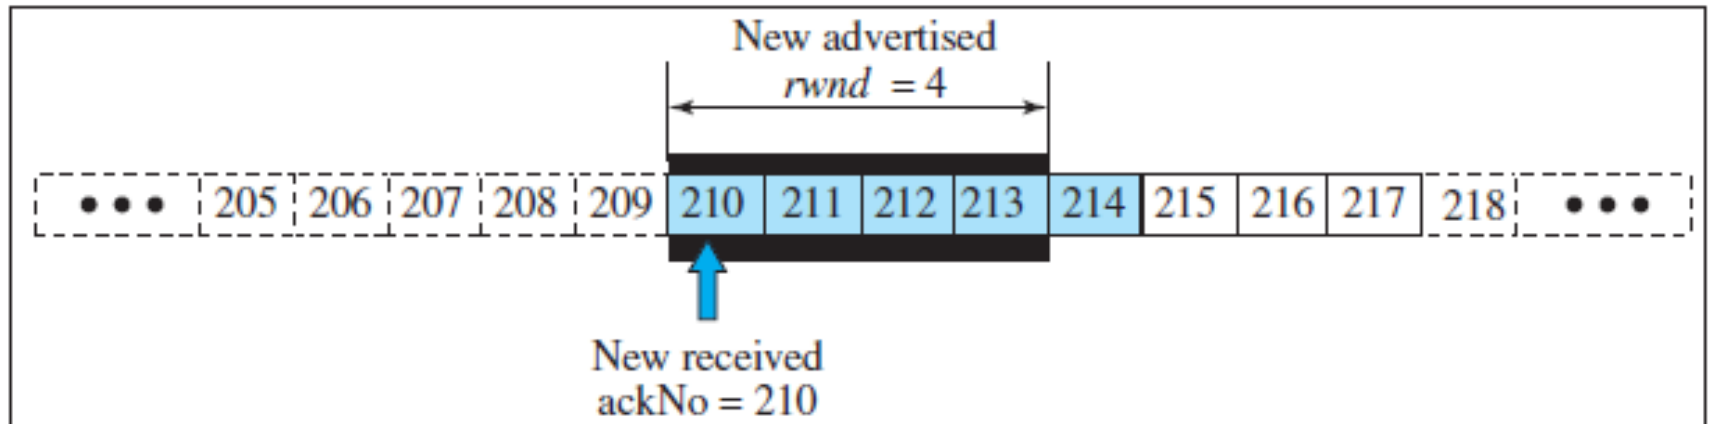
Data flow and flow control feedbacks in TCP



Shrinking of Window



a. The window after the last advertisement



b. The window after the new advertisement; window has shrunk

Window Shutdown

Silly Window Syndrome

- The sending application program creates data slowly or the receiving application program consumes data slowly, or both.
 - Creates problem in sliding window operation and leads to overhead.
- Syndrome created by
 - Sender
 - Receiver

Syndrome created by the Sender

- The sending TCP may create a silly window syndrome if it is serving an application program that creates data slowly
- Solution: prevent the sending TCP from sending the data in small quantity
 - Nagle's Algorithm
 - Send first piece of data received from sending app.
 - Accumulate data and wait until an ACK is received or until max data is accumulated

Syndrome created by the Receiver

- The receiving TCP may create a silly window syndrome if it is serving an application program that consumes data slowly
- Two Solutions:
 - send an ACK as soon as the data arrive, but to announce a window size of zero until either there is enough space to accommodate a segment of maximum size or until at least half of the receive buffer is empty (Clark's Solution)
 - delay sending the acknowledgment

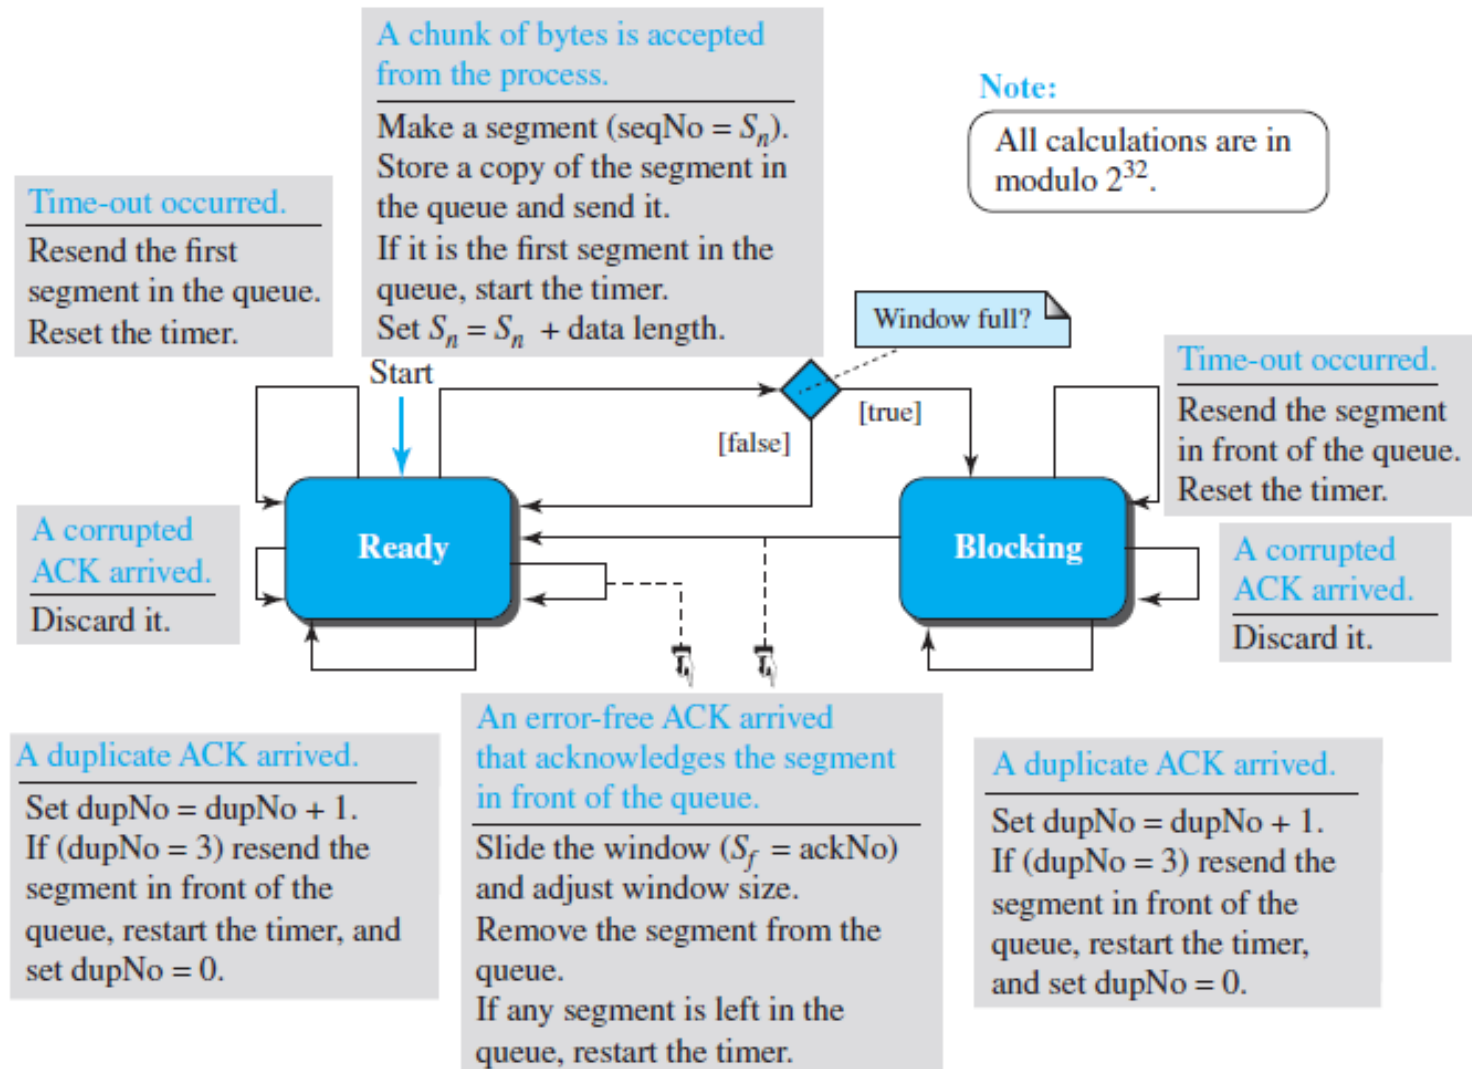
TCP – Error Control

- Detecting and resending corrupted segments
- Resending lost segments
- Storing out-of order segments until missing segments arrive
- Detecting and discarding duplicated segments
- Solution
 - checksum, acknowledgment (Cumulative, Selective), and time-out

TCP Acknowledgement

- Acknowledgement
 - Cumulative (ACK)
 - Selective (SACK)
- Retransmission
 - After RTO (Retransmission time-out)
 - After three duplicate ACK segments
 - Fast retransmission

FSM for the TCP sender side



FSM for the TCP receiver side

Note:

All calculations are in modulo 2^{32} .

An expected error-free segment arrived.

Buffer the message.

$R_n = R_n + \text{data length}$.

If the ACK-delaying timer is running, stop the timer and send a cumulative ACK.

Otherwise, start the ACK-delaying timer.

A request for delivery of k bytes of data from process came.


Deliver the data.

Slide the window and adjust window size.

An error-free duplicate segment or an error-free segment with sequence number outside window arrived.

Discard the segment.

Send an ACK with ackNo equal to the sequence number of expected segment (duplicate ACK).



ACK-delaying timer expired.

Send the delayed ACK.

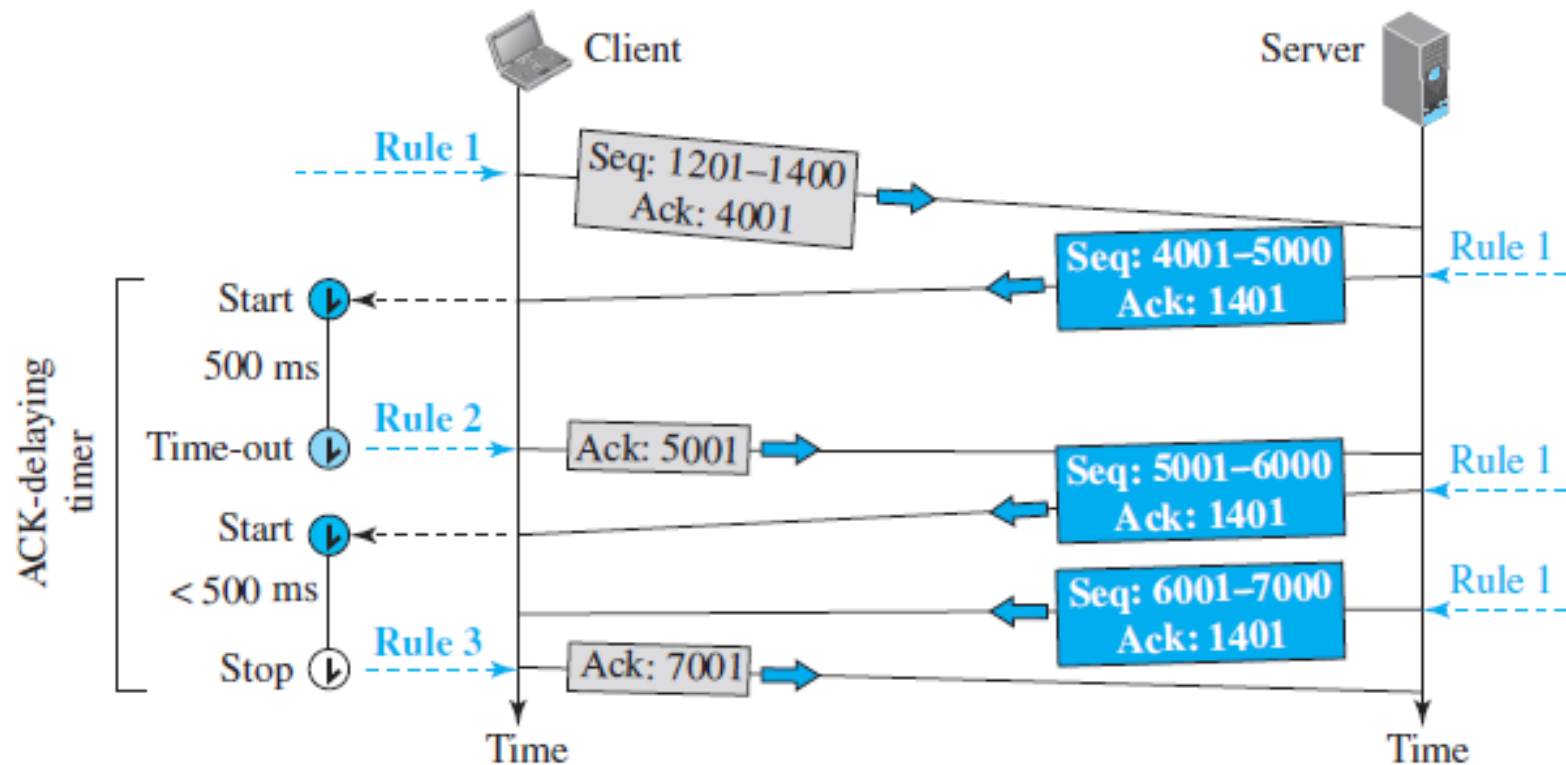
An error-free, but out-of-order segment arrived.

Store the segment if not duplicate. Send an ACK with ackNo equal to the sequence number of expected segment (duplicate ACK).

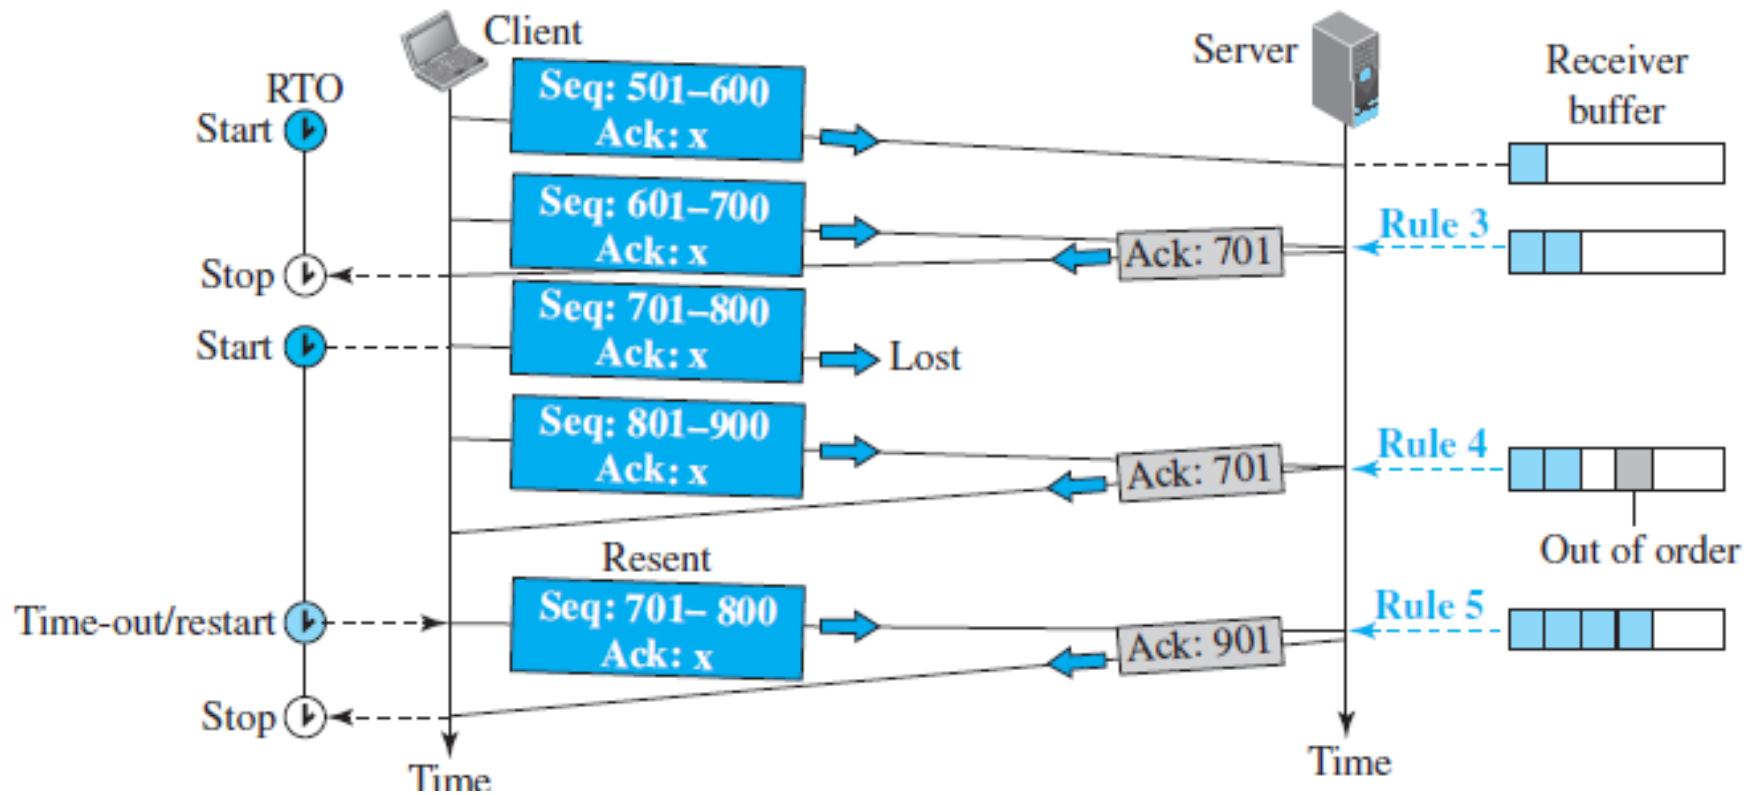
A corrupted segment arrived.

Discard the segment.

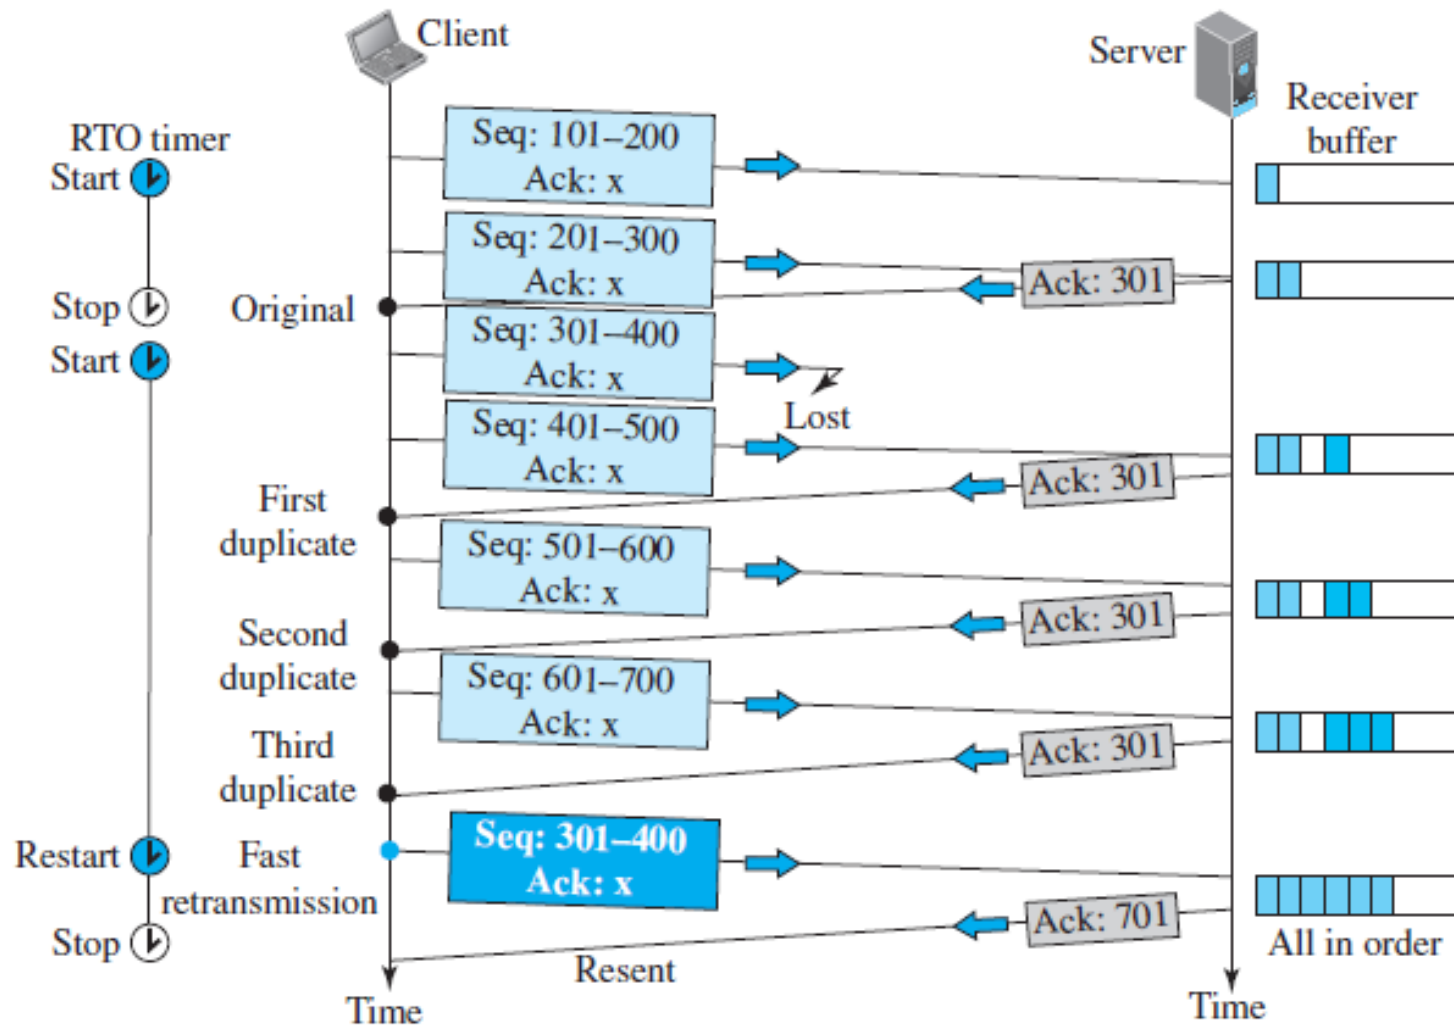
Normal Operations



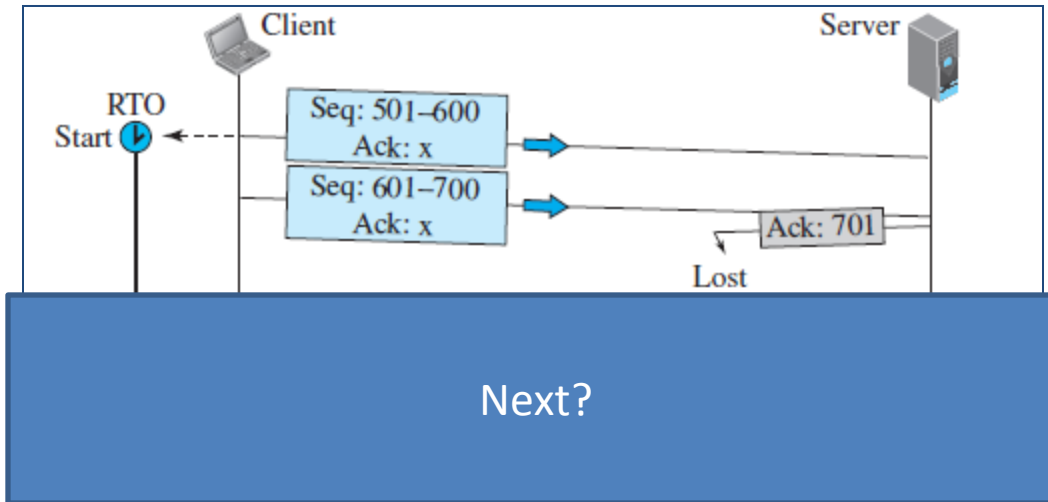
Lost Segment



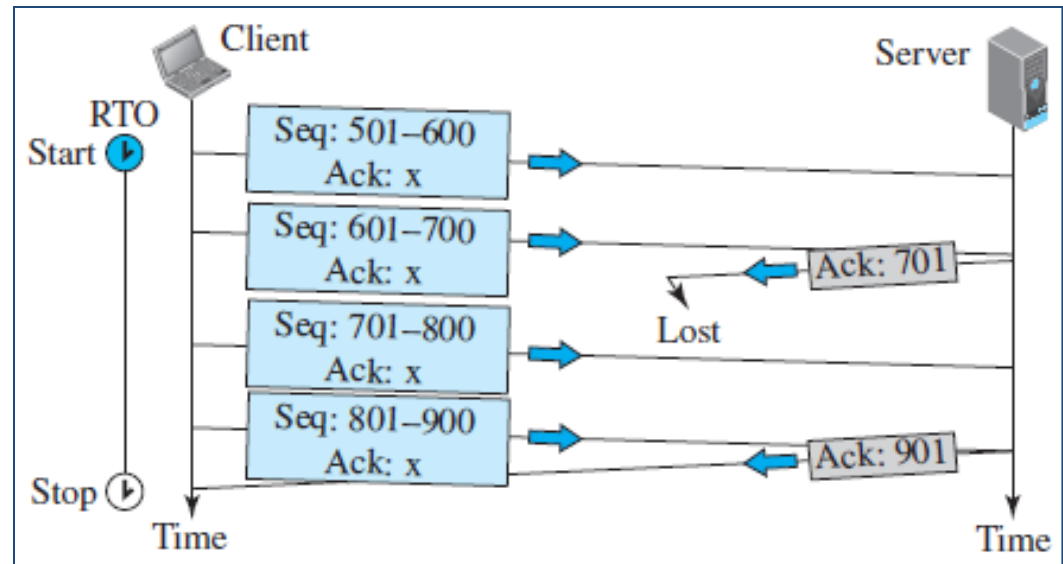
Fast retransmission



Lost acknowledgement



Deadlock due to Lost acknowledgement



TCP – Congestion Control

- Implemented at the **Sender** side
- Send window = $\text{MIN}(rwnd, cwnd)$
 - *rwnd*: advertised by the receiver
 - *cwnd*: adjusted based on feedback from the network
- Congestion Detection
 - **Time-out** – Strong congestion
 - **Three duplicate ACKs** – Weak congestion

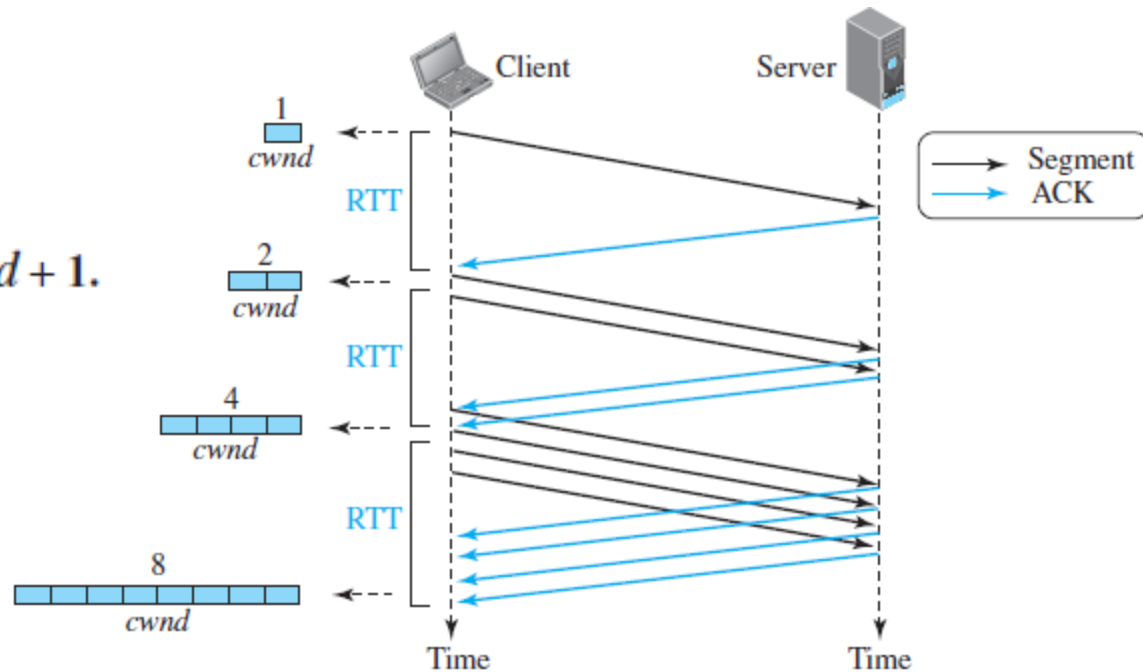
Congestion Policies

- Slow start: Exponential increase
- Congestion avoidance
- Fast recovery

Slow Start

Initially, $cwnd = 1$ MSS

If an ACK arrives, $cwnd = cwnd + 1$.



Start

After 1 RTT

After 2 RTT

After 3 RTT

→ $cwnd = 1 \rightarrow 2^0$

→ $cwnd = cwnd + 1 = 1 + 1 = 2 \rightarrow 2^1$

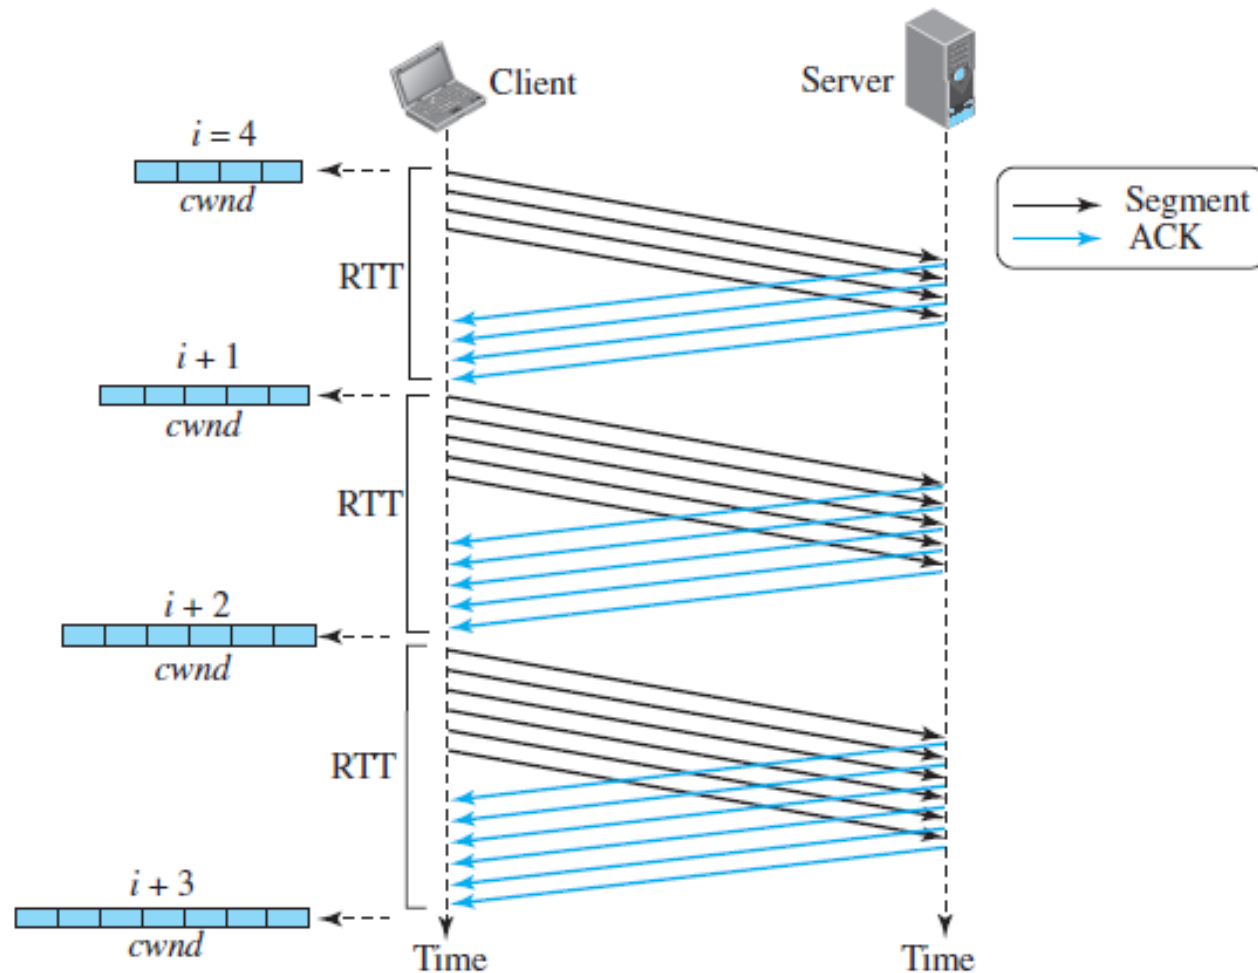
→ $cwnd = cwnd + 2 = 2 + 2 = 4 \rightarrow 2^2$

→ $cwnd = cwnd + 4 = 4 + 4 = 8 \rightarrow 2^3$

Congestion Avoidance

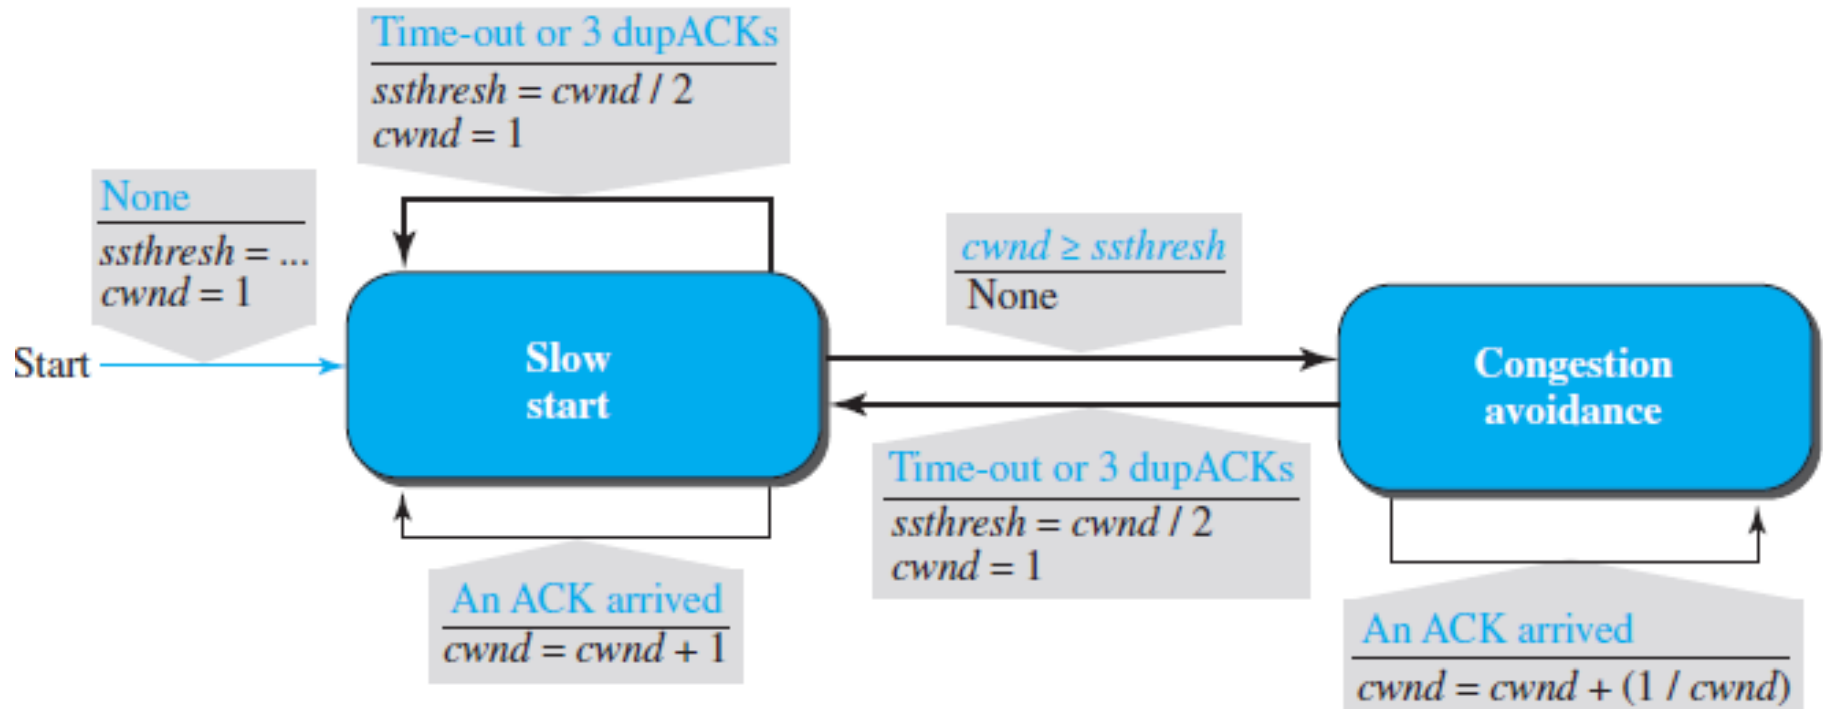
- Congestion avoidance phase is started if *cwnd* has reached the *slow-start threshold* value
- If ***cwnd* \geq *ssthresh*** then each time an ACK is received, increment *cwnd* as follows:
 - $cwnd = cwnd + 1/cwnd$
- So *cwnd* is increased by one segment (=MSS bytes) only if all segments have been acknowledged.

Congestion Avoidance




TCP – Congestion Control

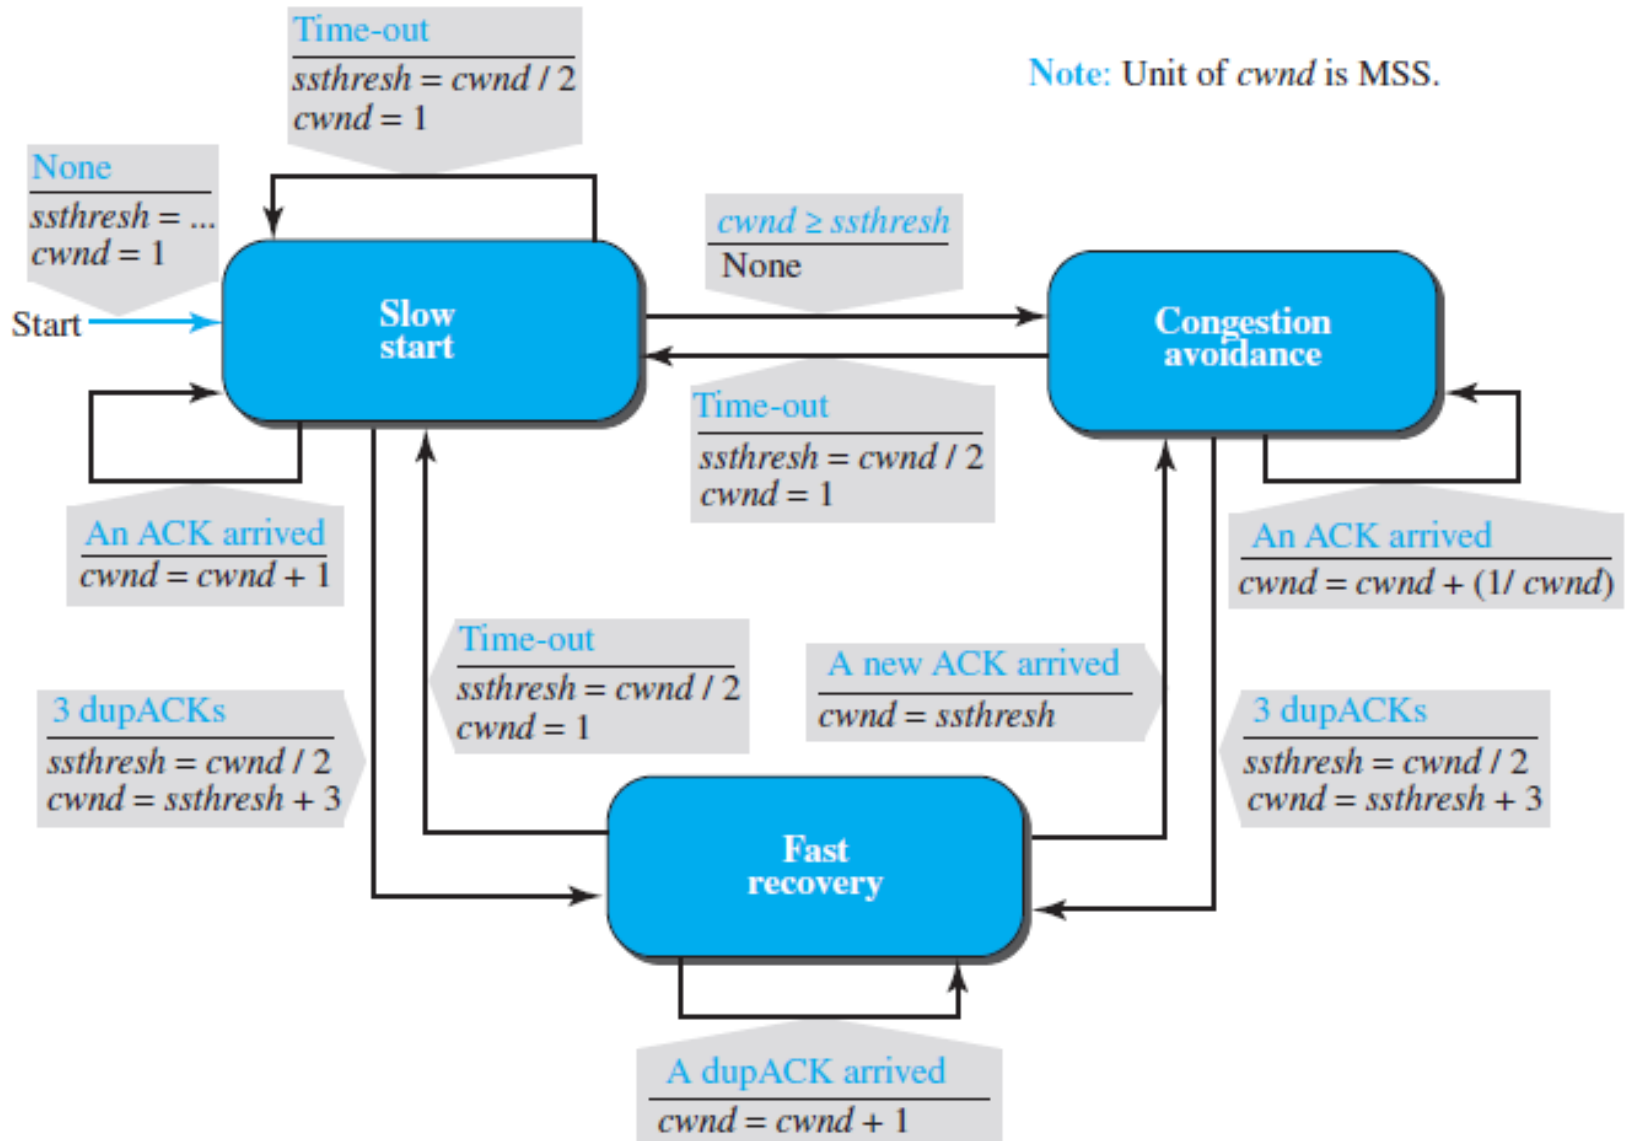
- Tahoe TCP
 - Slow start and Congestion Avoidance




Tahoe TCP – Example



Reno TCP



Reno TCP – Example



AIMD

- Additive increase, multiplicative decrease (AIMD)

